Problem-First AI

Lecture 0:

Foundations and choosing the right problems

Aishwarya Reganti

Kiriti Badam

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What we'll cover

- Applied AI this year: What has changed?
- How do AI models work?
- How to choose the right problems for AI?
- System Design for AI

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

01 Applied AI this year

What has changed, and which kinds of work can we now tackle?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

How applied AI has evolved

	2023	2024	2025	Q4 2026
Consumer	**Chatbots** 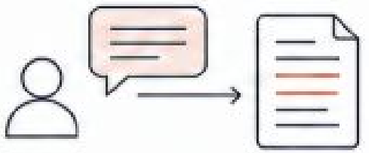 Ask, explain, draft	**Assistants** 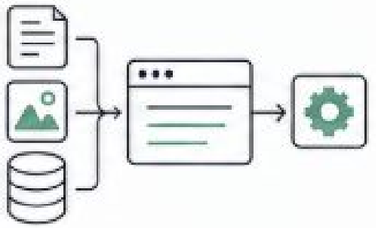 Read files and use tools	**Teammates** 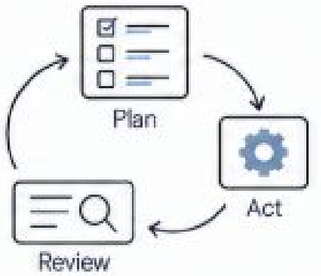 Delegate a longer assignment	**Autonomous workforce** 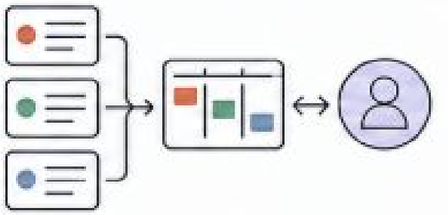 Coordinate agents; approve actions
Enterprise	**Knowledge access** Internal Q&A and drafting	**Business integration** 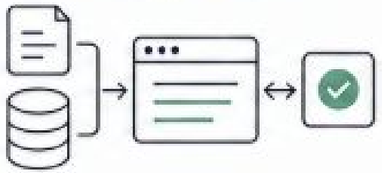 CRM queries and record updates	**Workflow execution** 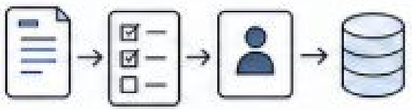 Invoice matching and approval routing	**End-to-end process transformation** 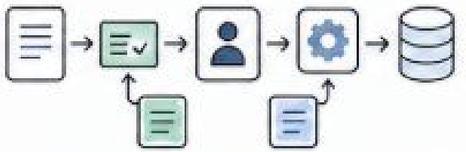 Procurement, from request to payment

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

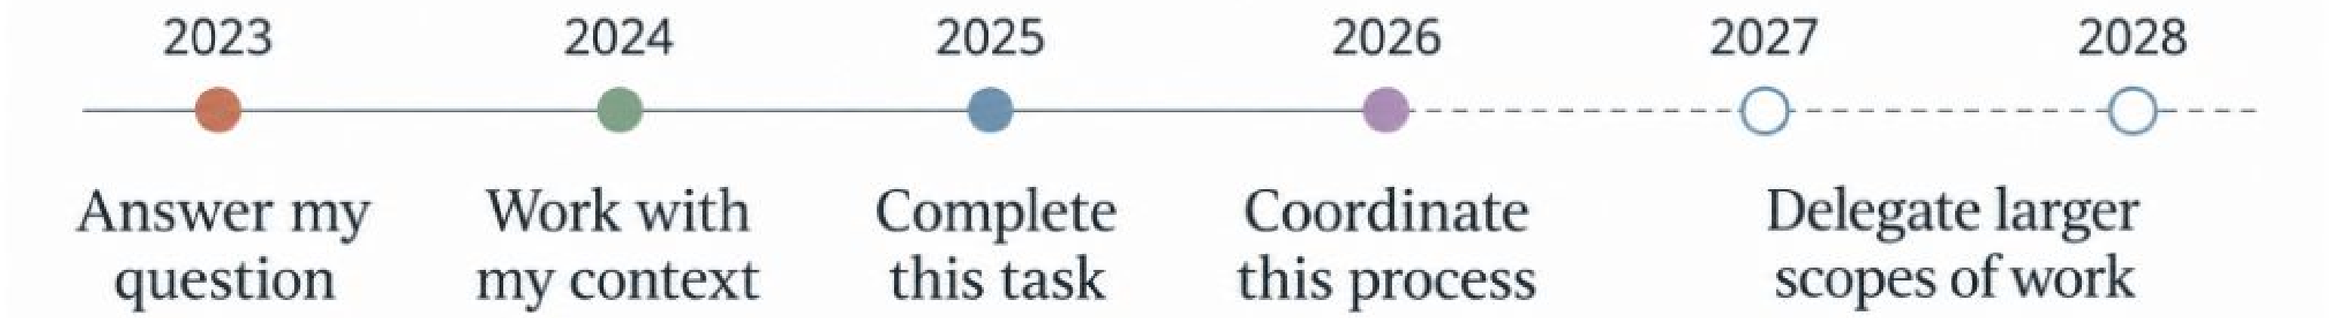

- We’re delegating larger pieces of work, from individual questions to tasks and whole processes.
- We expect AI to work over longer periods, keeping track of progress across more steps.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Horizontal, departmental, and vertical AI

Horizontal

- Used across many roles
- Example: Muse across business tools

Departmental

- Built for a business function
- Example: finance reconciliation

Vertical

- Built for an industry's workflows
- Example: healthcare documentation

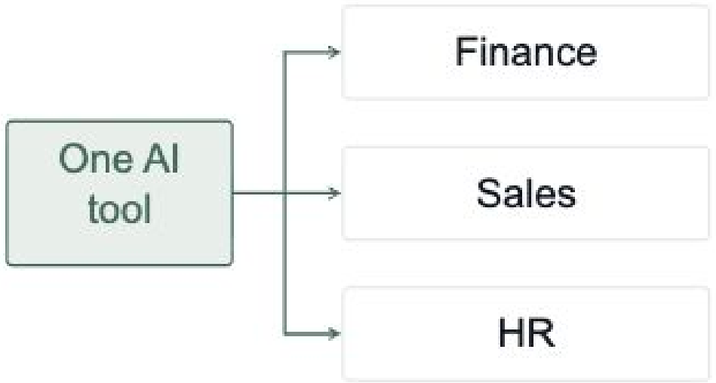

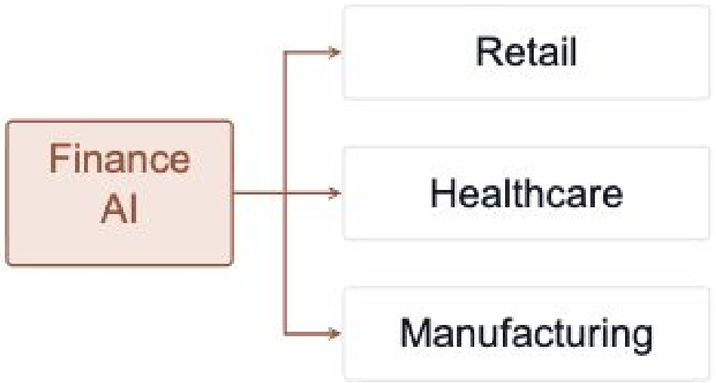

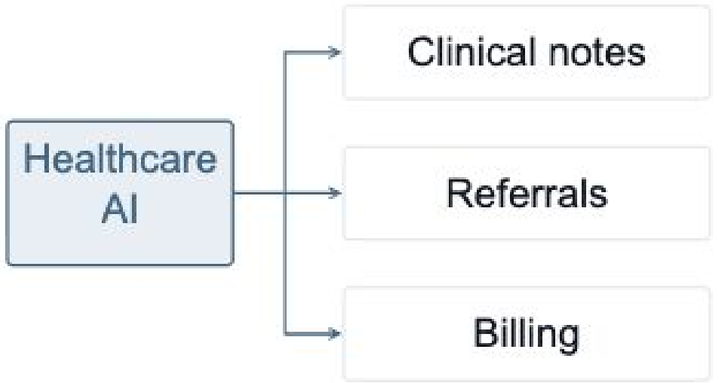

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Departmental AI Market Map

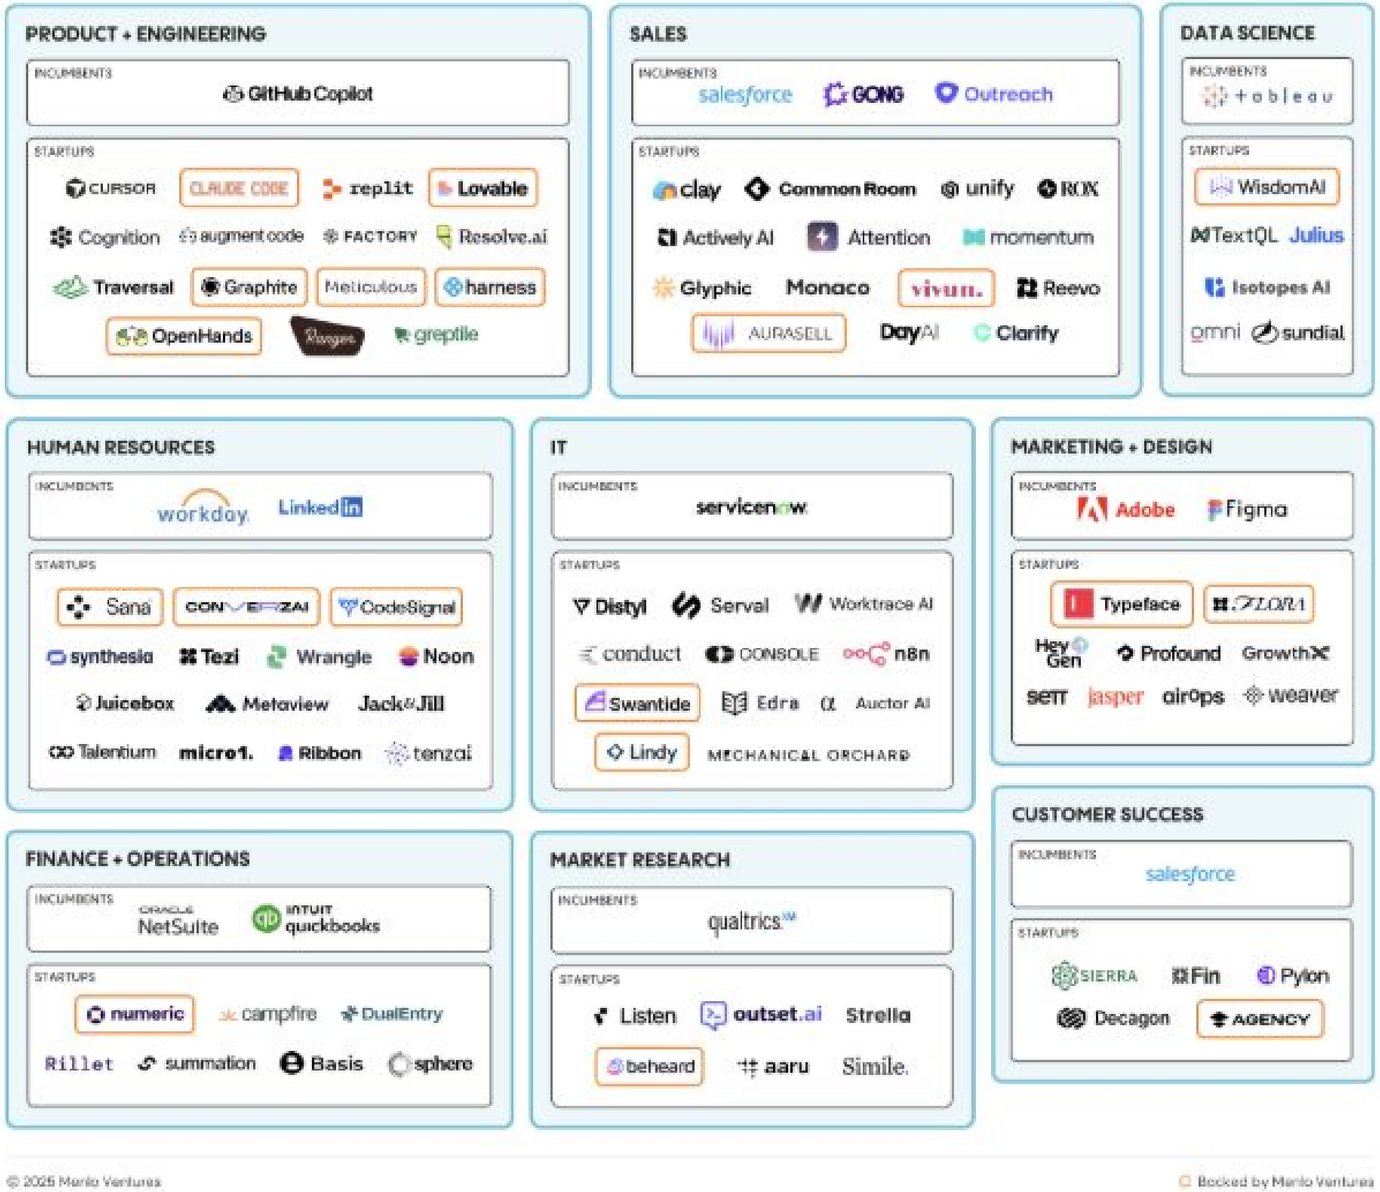

Source: Menlo Ventures, Departmental AI Market Map (2025)

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Vertical AI Market Map

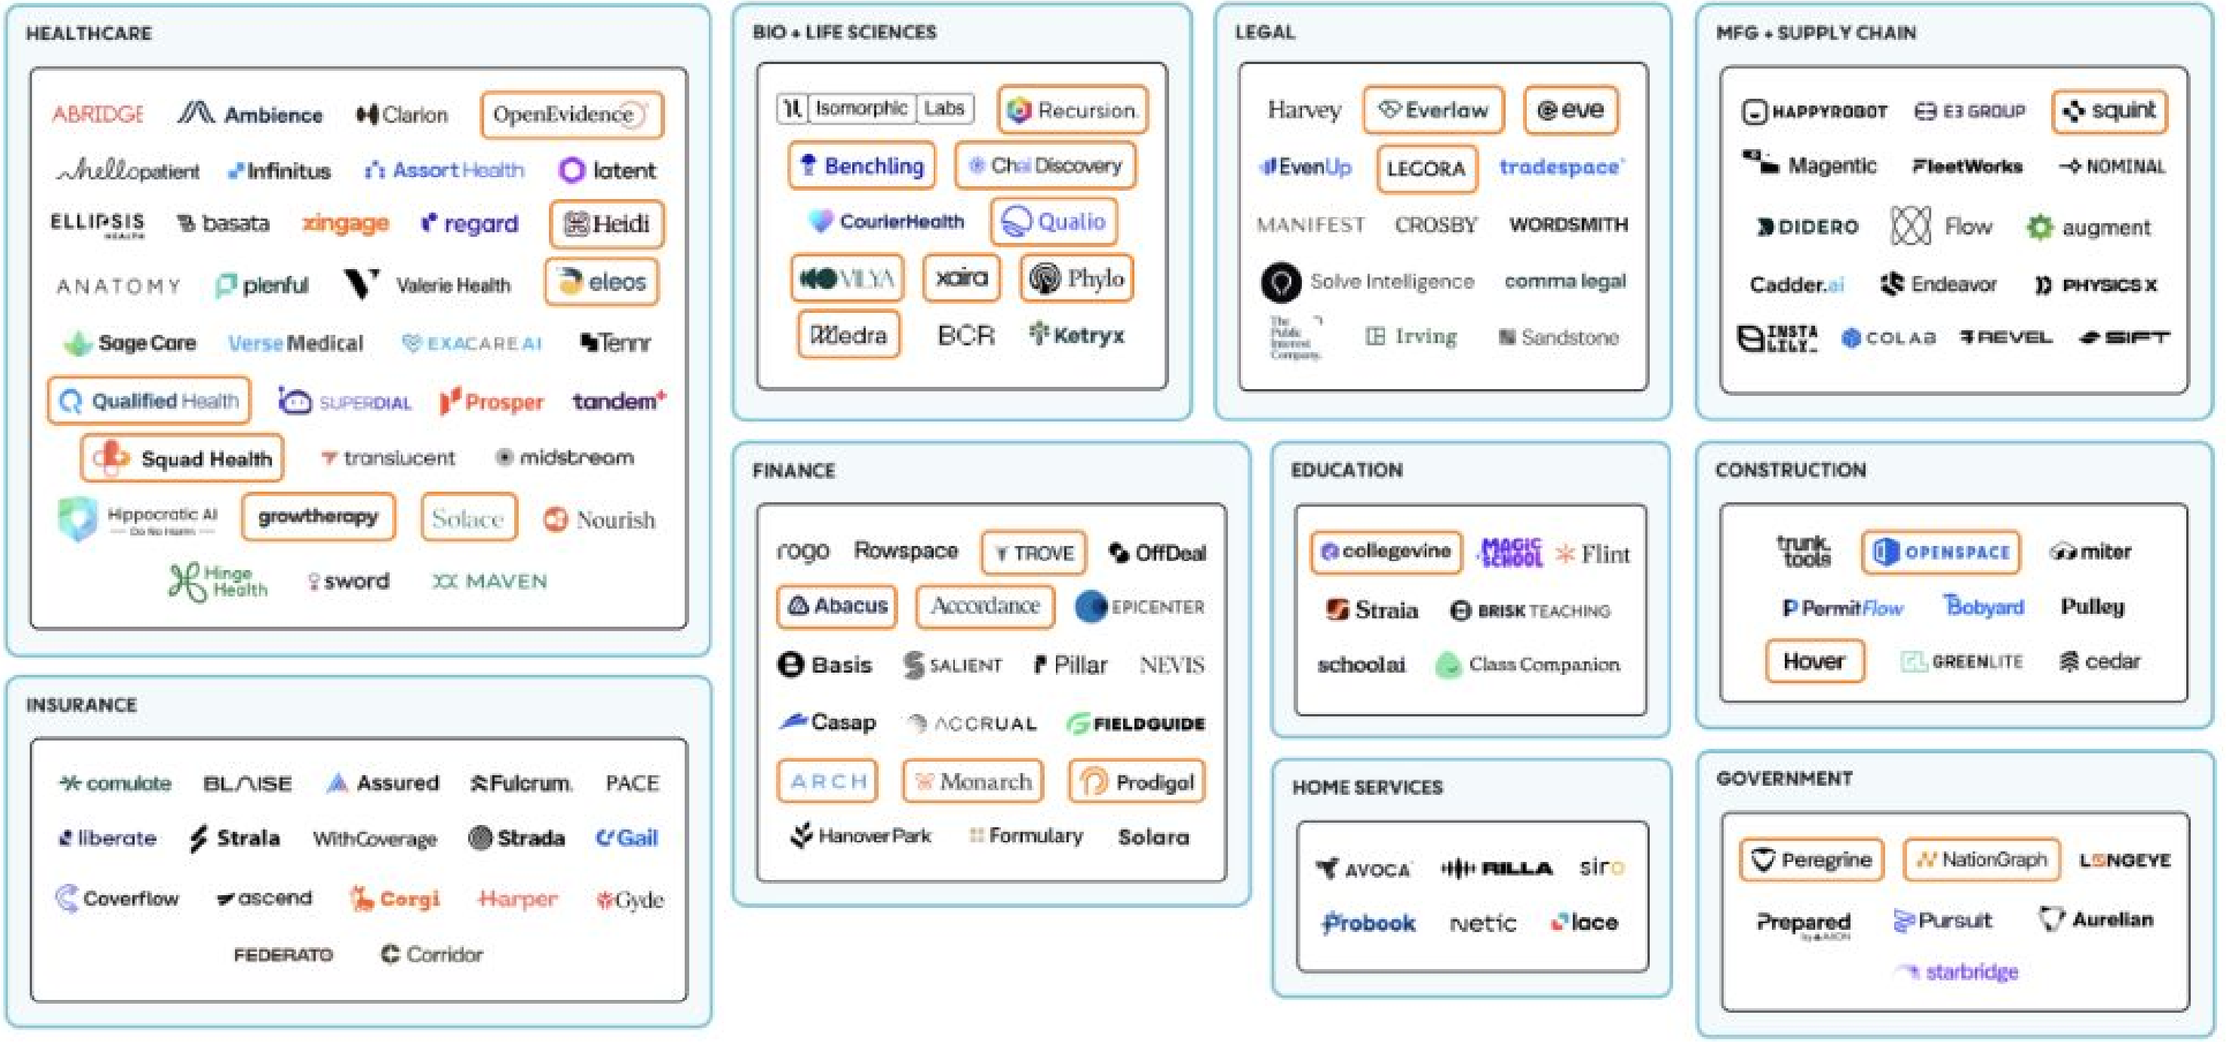

© 2026 Menlo Ventures

Menlo Ventures portfolio company

Source: Menlo Ventures, Vertical AI Market Map (2026)

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Why specialization creates value

- **Broader scope demands deeper knowledge.**
  Completing work requires the rules, exceptions and judgment of a department or industry.

- **Value depends on fitting the workflow.**
  AI needs the right records, systems and approvals to deliver a usable result.

- **Depth creates room to differentiate.**
  Specialized integrations, quality checks and feedback are harder to replicate than a generic AI interface.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

We’re bullish on departmental and vertical AI.

- AI can take on larger pieces of work, opening up more business processes to software.
- Delivering that work requires deeper knowledge of the department or industry.
- For builders, that depth is where we see the opportunity to create value.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

LevelUp Lens™

Departmental and vertical AI will create huge opportunities for builders within enterprises as well as builders looking to build successful products in 2026 and beyond.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

02 How do AI models work?

We've seen the 10,000-foot view of applied AI. Now let's look inside the models that make these products possible.

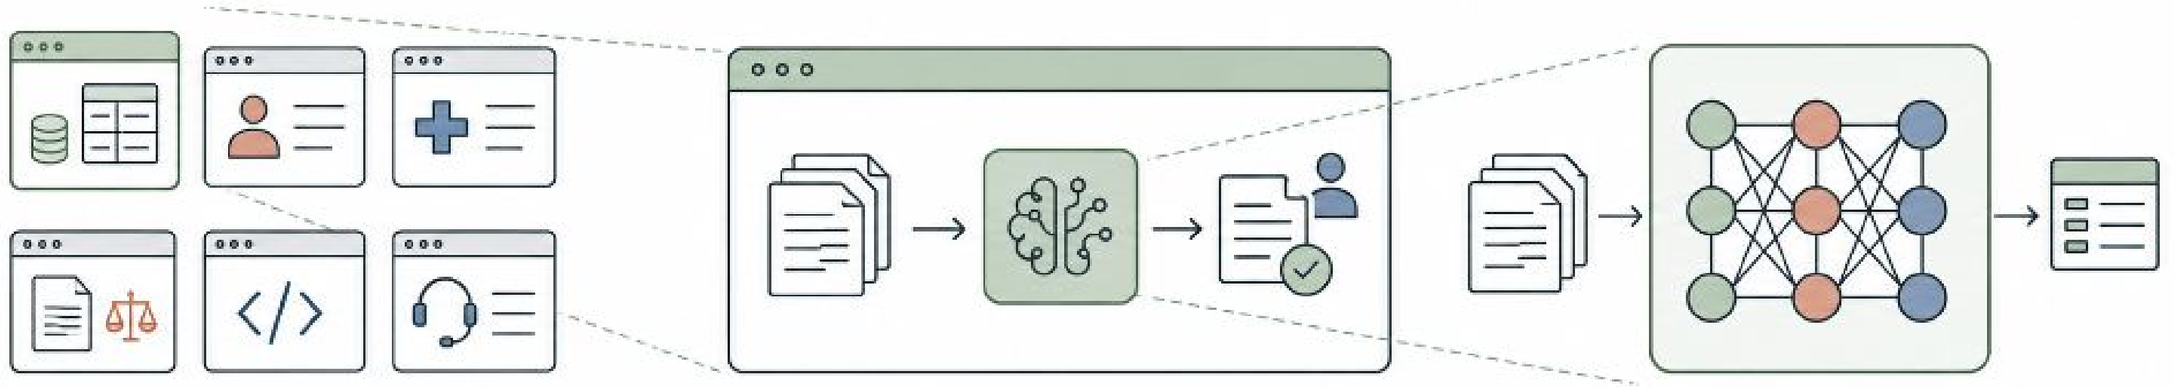

AI product landscape | One business workflow | The model inside it

To build useful AI applications, we need to understand what models learn, what they need as input, and where they fall short.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What is artificial intelligence?

AI is the field of computer science focused on creating systems that perform tasks that typically require human intelligence.

Aa Language

Patterns

? Decisions

Content

These include understanding language, recognizing patterns, making decisions, solving problems and generating content.

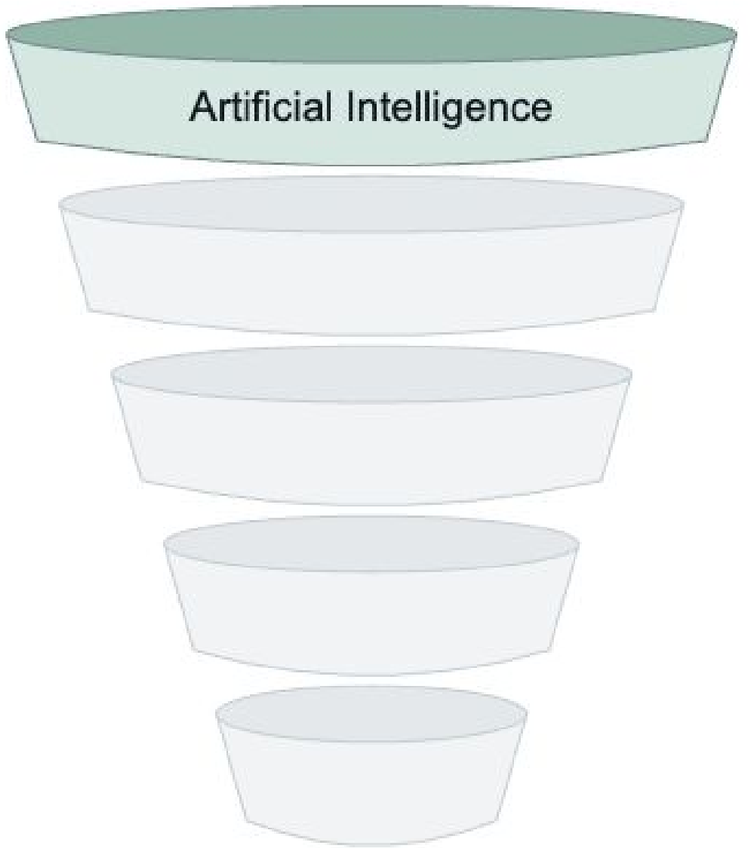

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What is machine learning?

Machine learning lets systems learn patterns from data instead of explicitly programming every rule.

→ **ML model**
Learned parameters
→ Dog

A model is a mathematical framework that turns inputs into outputs using learned patterns.

Its parameters are adjusted during training.

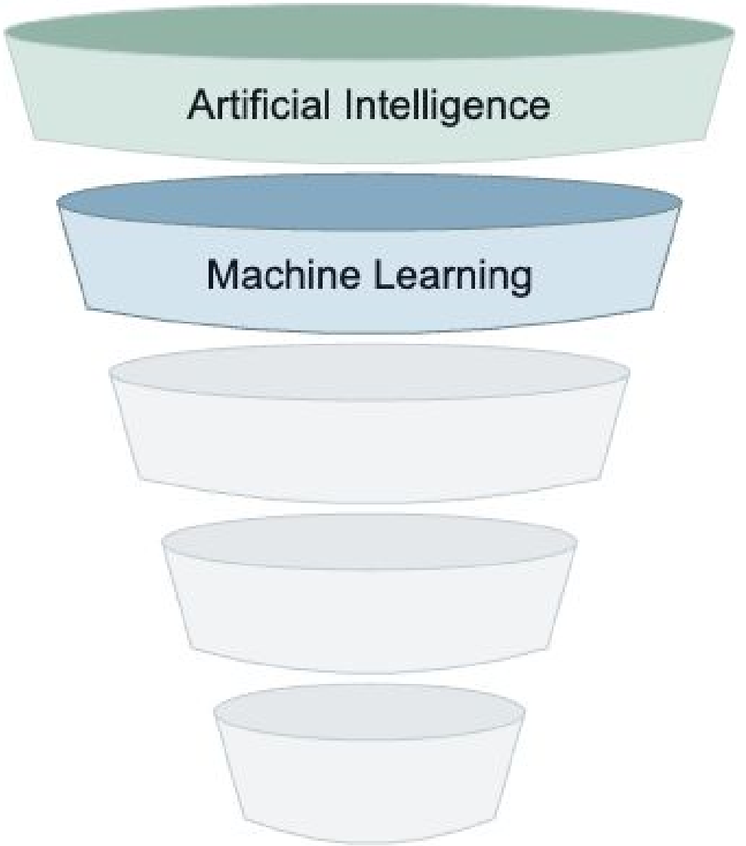

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What are model parameters?

Parameters are numbers inside a model that are adjusted during training to change how it turns inputs into outputs.

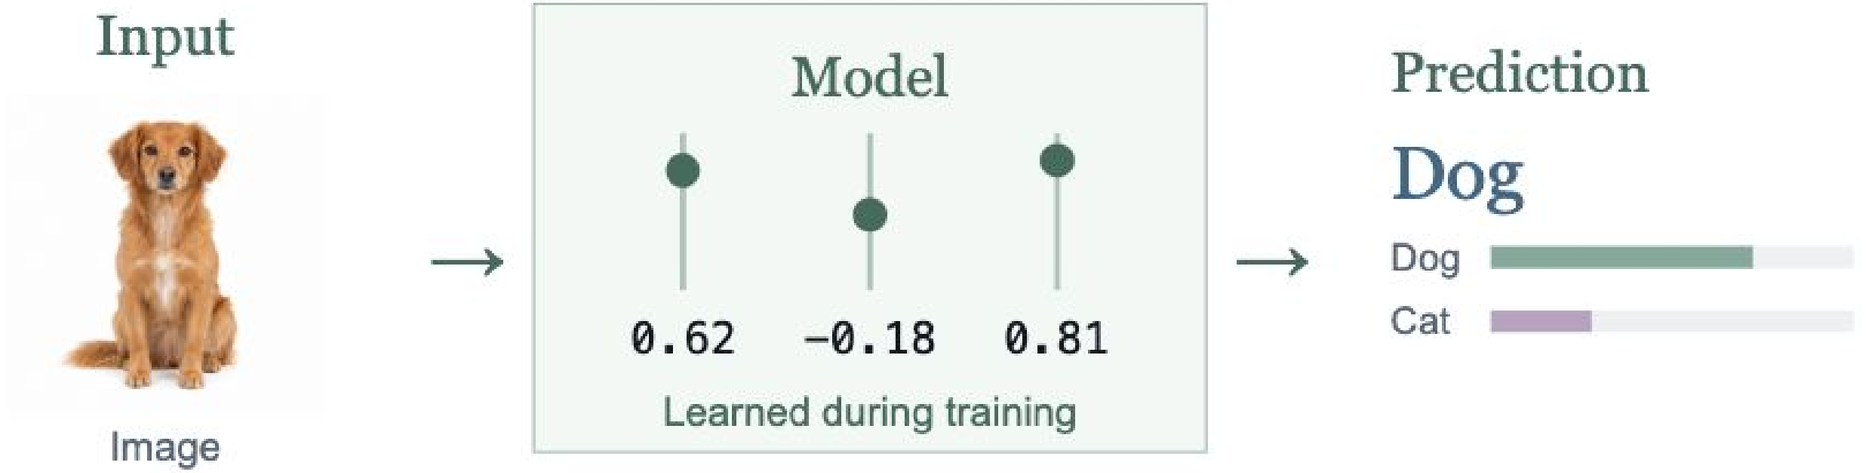

Training is how the model's parameters are adjusted using data and a learning signal.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Training: learning from examples

More examples

1

Predict

Use current parameters

→

2

Compare

Use the correct label

→

3

Adjust

Update parameters

Repeat across many examples, using the updated parameters each time.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Supervised learning

The examples come with the answers we want the model to learn.

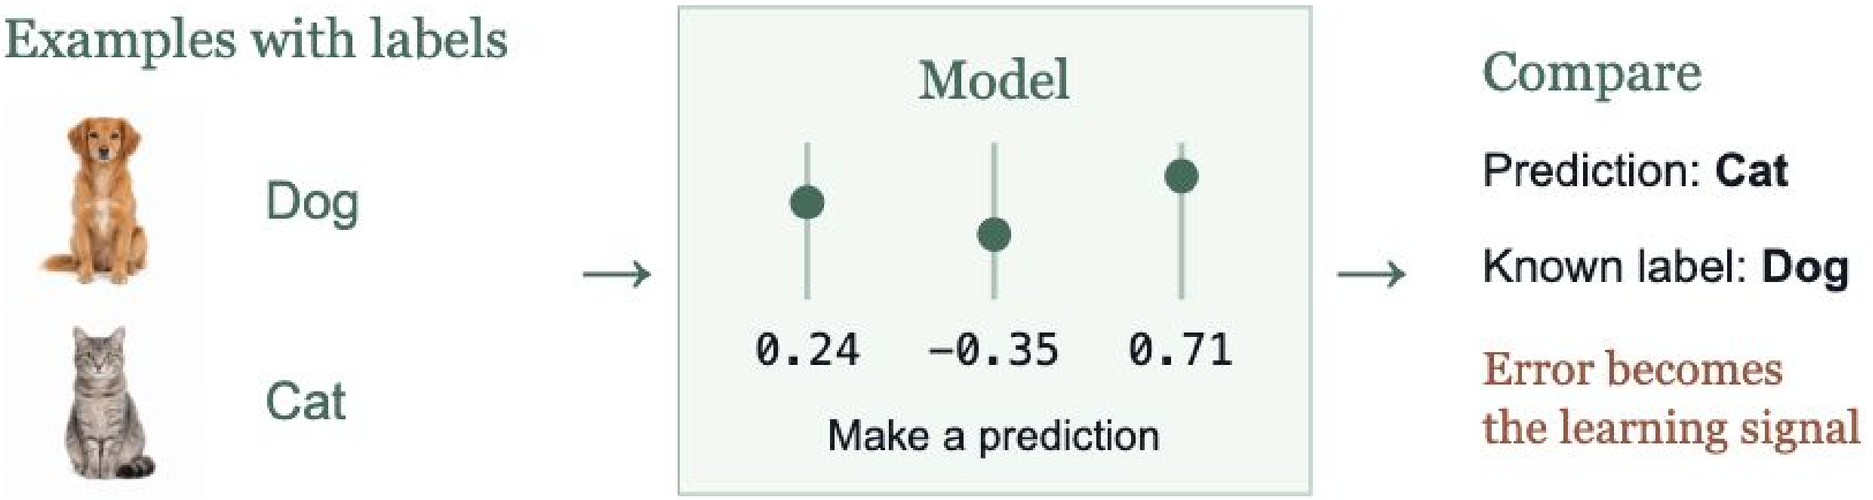

Use the error as feedback to adjust parameters. Repeat across examples.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Unsupervised learning

Find similar characteristics in images, without being given “dog” or “cat” labels.

Dogs

Rounded ears · longer muzzle

Cats

Pointed ears · shorter muzzle

A person can then look at the groups and name them dogs and cats.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Self-supervised learning

Create a prediction task from the data itself.

Hide a piece

→

Predict it

→

Check the original

Use the prediction error to adjust the model.

Repeat across images. Each image supplies both the input and the answer.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Inference: using the trained model

Inference uses learned parameters to produce an output for the current input.

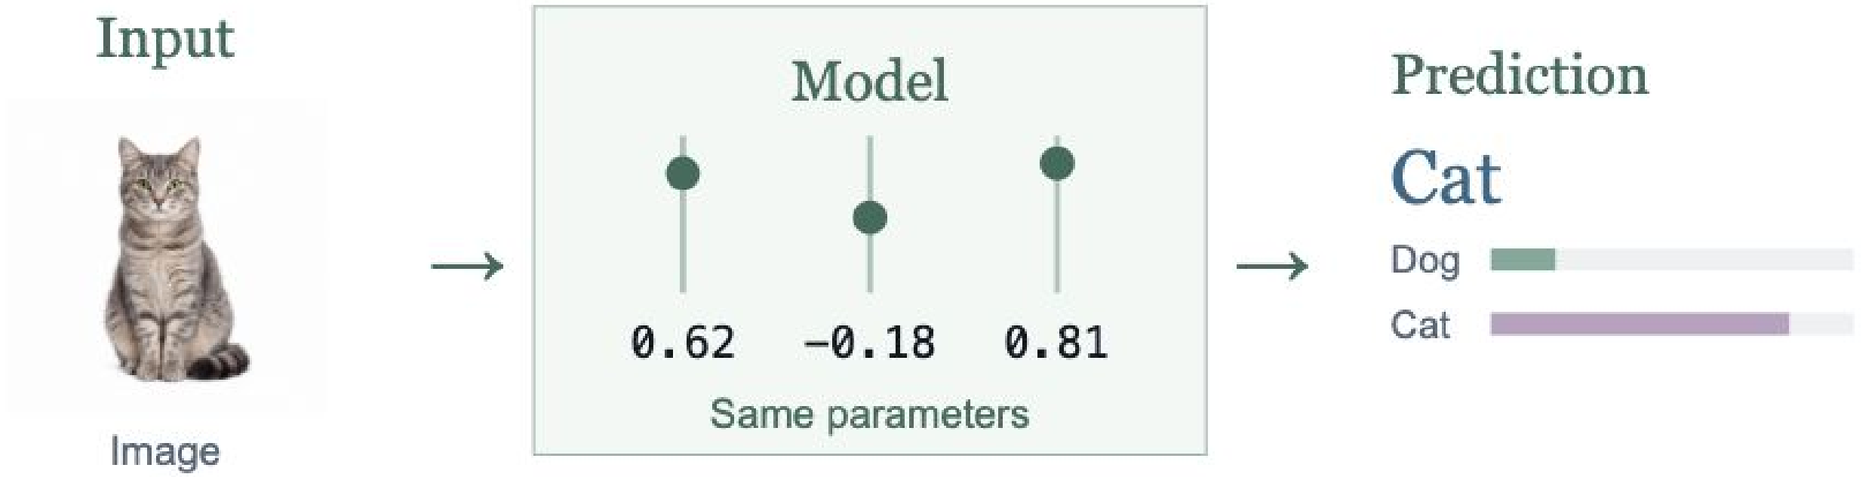

Using the model is inference. Updating its parameters is training.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Training and inference

The same model plays two different roles.

Training

Examples + learning signal

↓

Adjust parameters

↻

Repeat across examples

Changes the model's learned settings.

Inference

Current input

↓

Use learned parameters

↓

Produce an output

Uses those settings without updating them.

An application can save information and supply it next time as context, without retraining the model.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What are neural networks?

A neural network is a machine learning model made of connected layers of nodes.

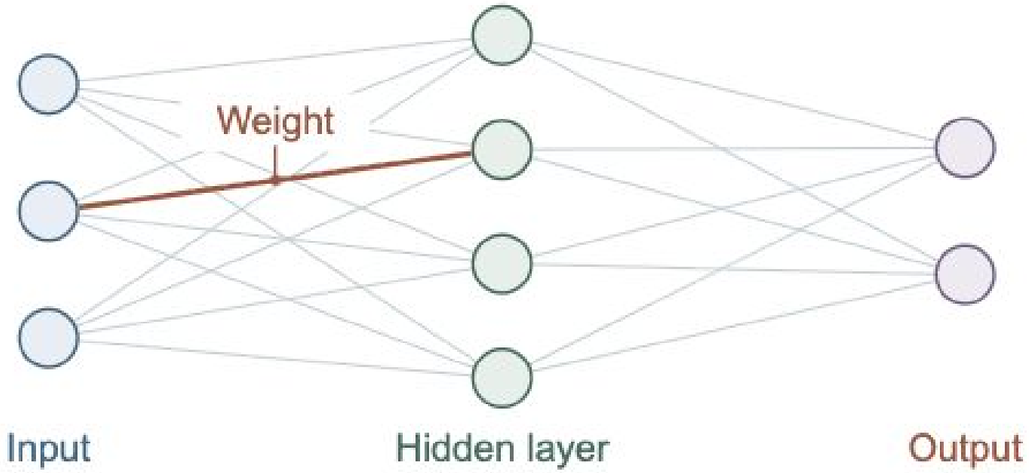

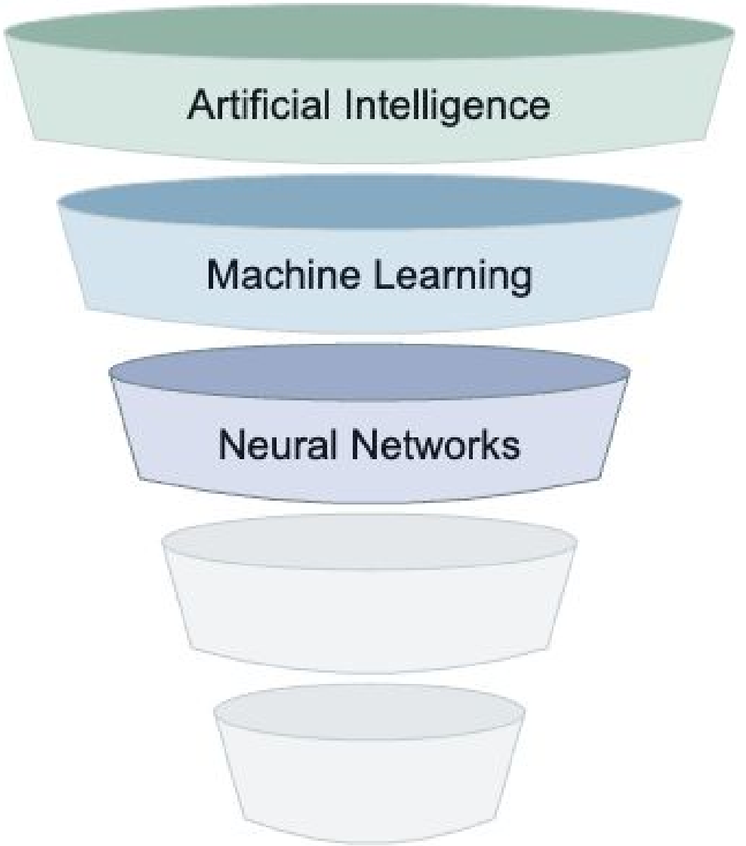

Inputs pass through weighted connections. Hidden layers transform them into an output.

Training changes the weights and biases.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Inside a neuron: weights and bias

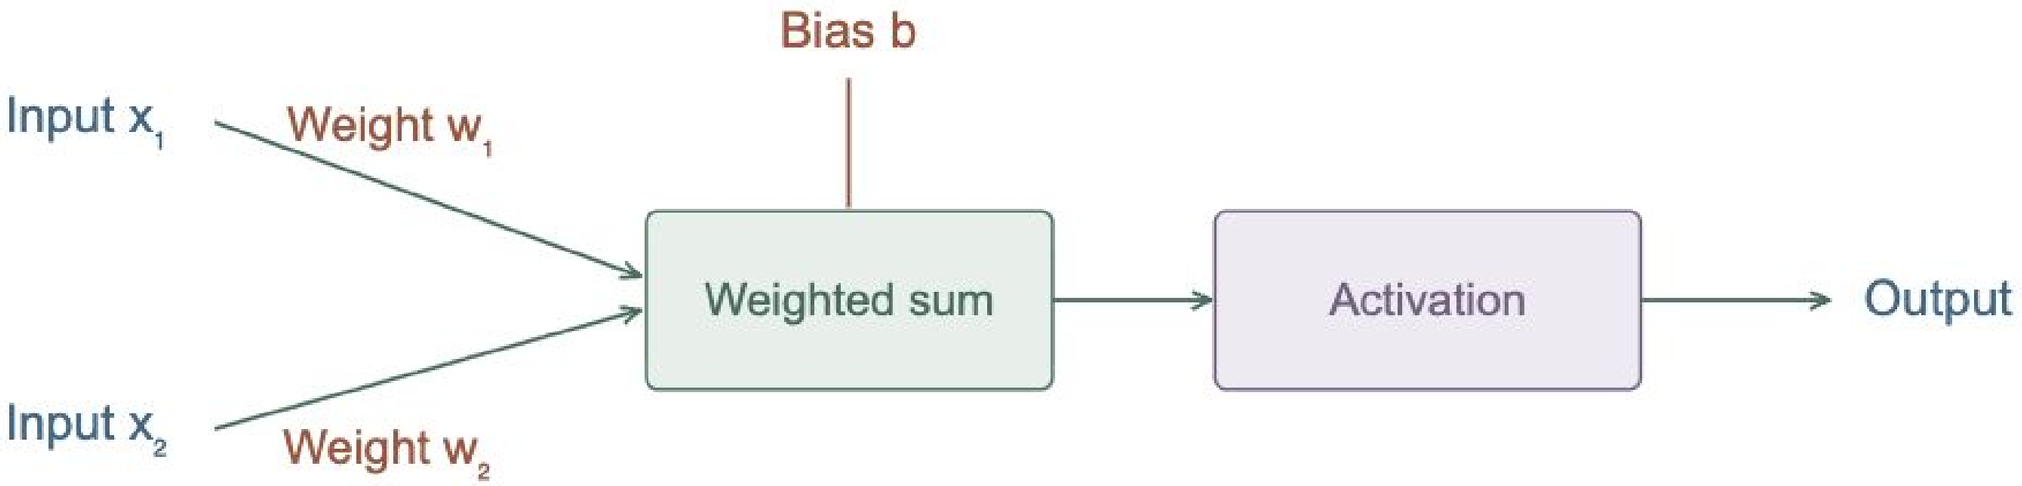

$$z = w_1x_1 + w_2x_2 + b \quad \rightarrow \quad \text{output} = f(z)$$

3 learned parameters 2 weights + 1 bias

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What is deep learning?

Deep learning uses neural networks with multiple learned layers to recognize patterns and make predictions.

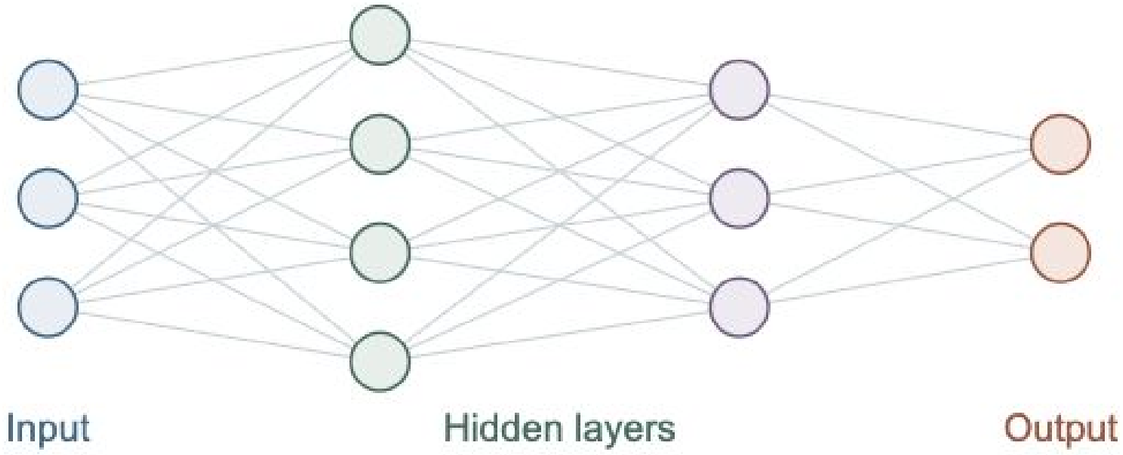

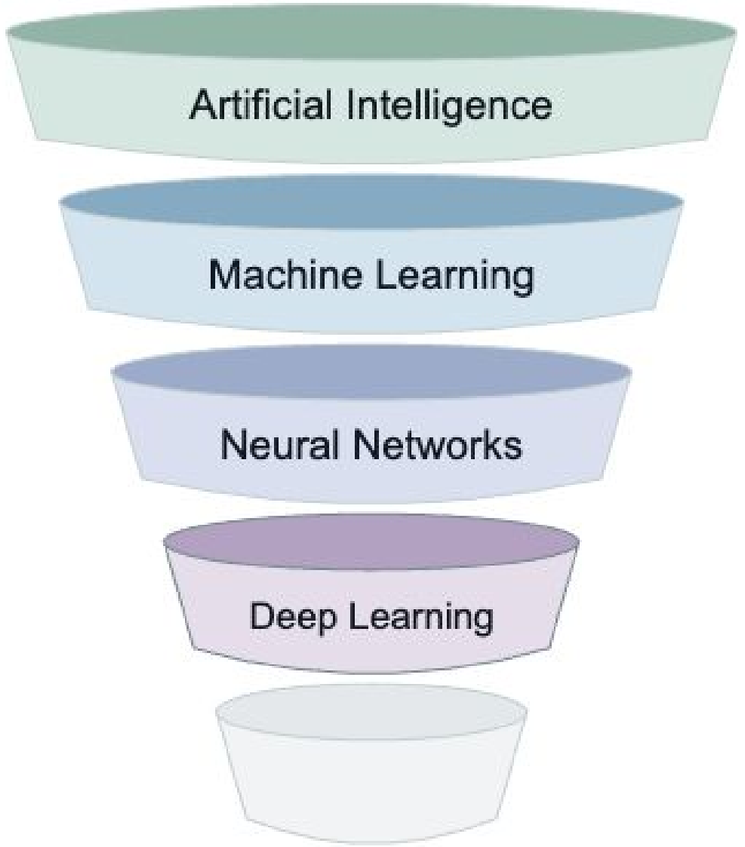

Each layer transforms the representation produced by the layer before it.

The hidden layers are learned, not hand-written rules.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Different neural-network architectures

Convolution

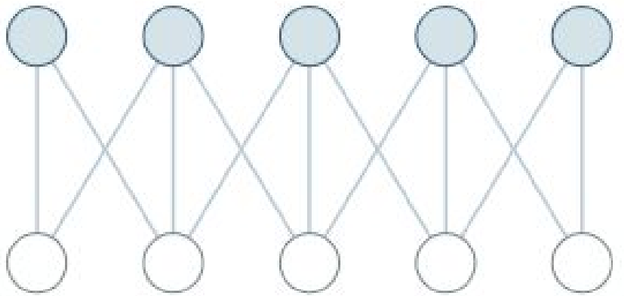

Connect local neighborhoods.

Recurrence

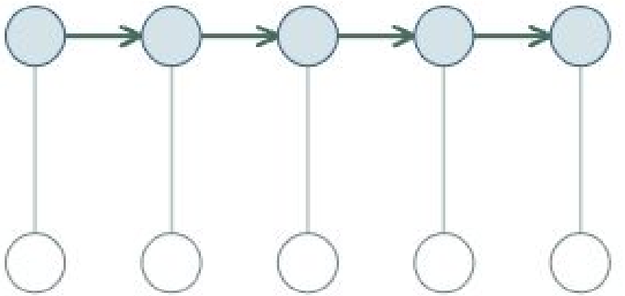

Carry state across a sequence.

Attention

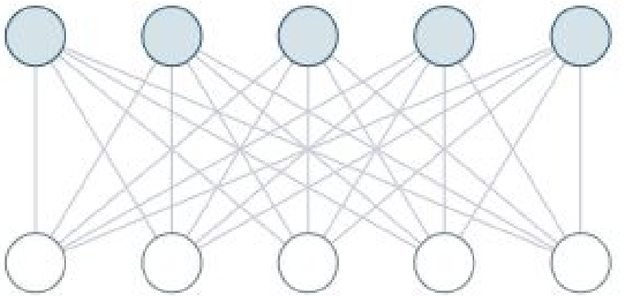

Combine information across positions.

A Transformer block Attention → Feed-forward network → Next block

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What is generative AI?

Generative AI models produce new content, such as text, images or audio.

Classify

→ Dog or cat

Choose from a defined set of labels.

Generate

"Describe this dog" → A golden-brown dog with a white chest…

Generate a sequence from many possible outputs.

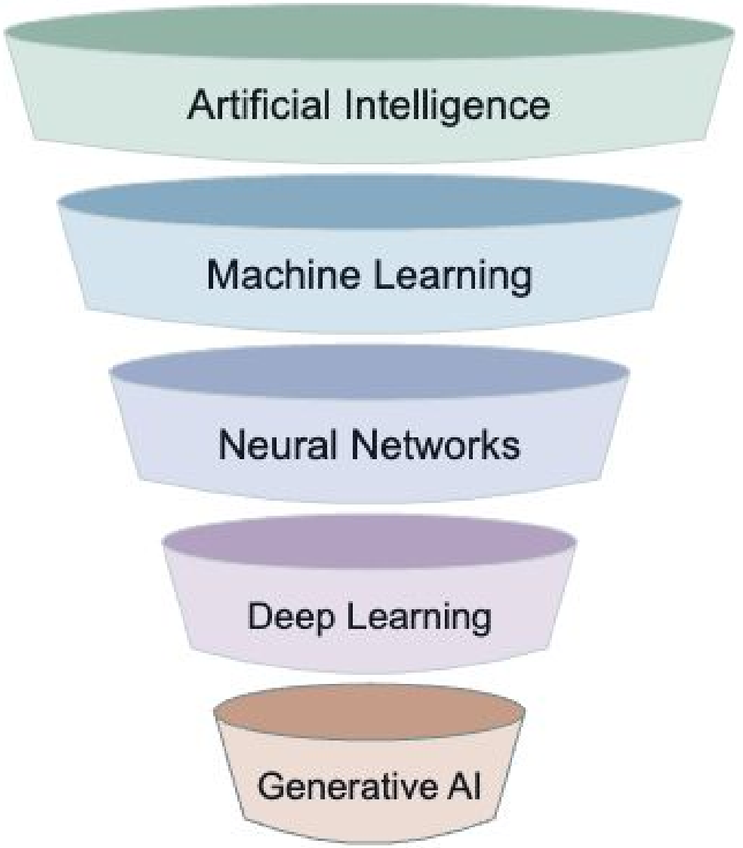

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What is agentic AI?

Agentic AI is an application layer built on top of generative AI models.

We add components that let the system act on what the model generates.

What do we add to turn a model into an agent?

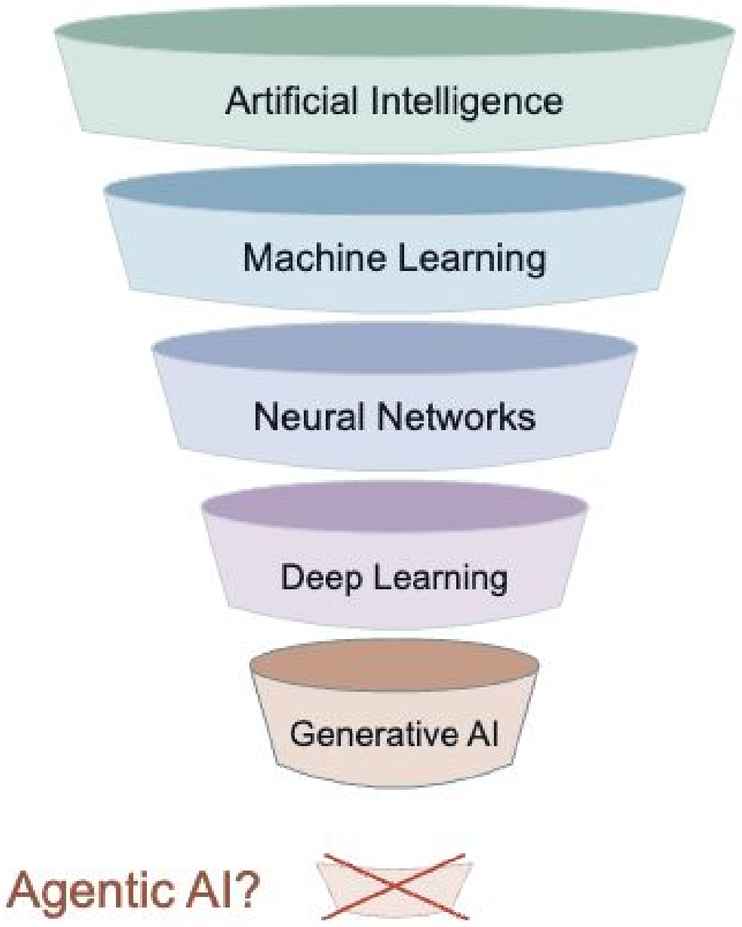

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Components of agentic AI

Add components that let the system act on the model's generations.

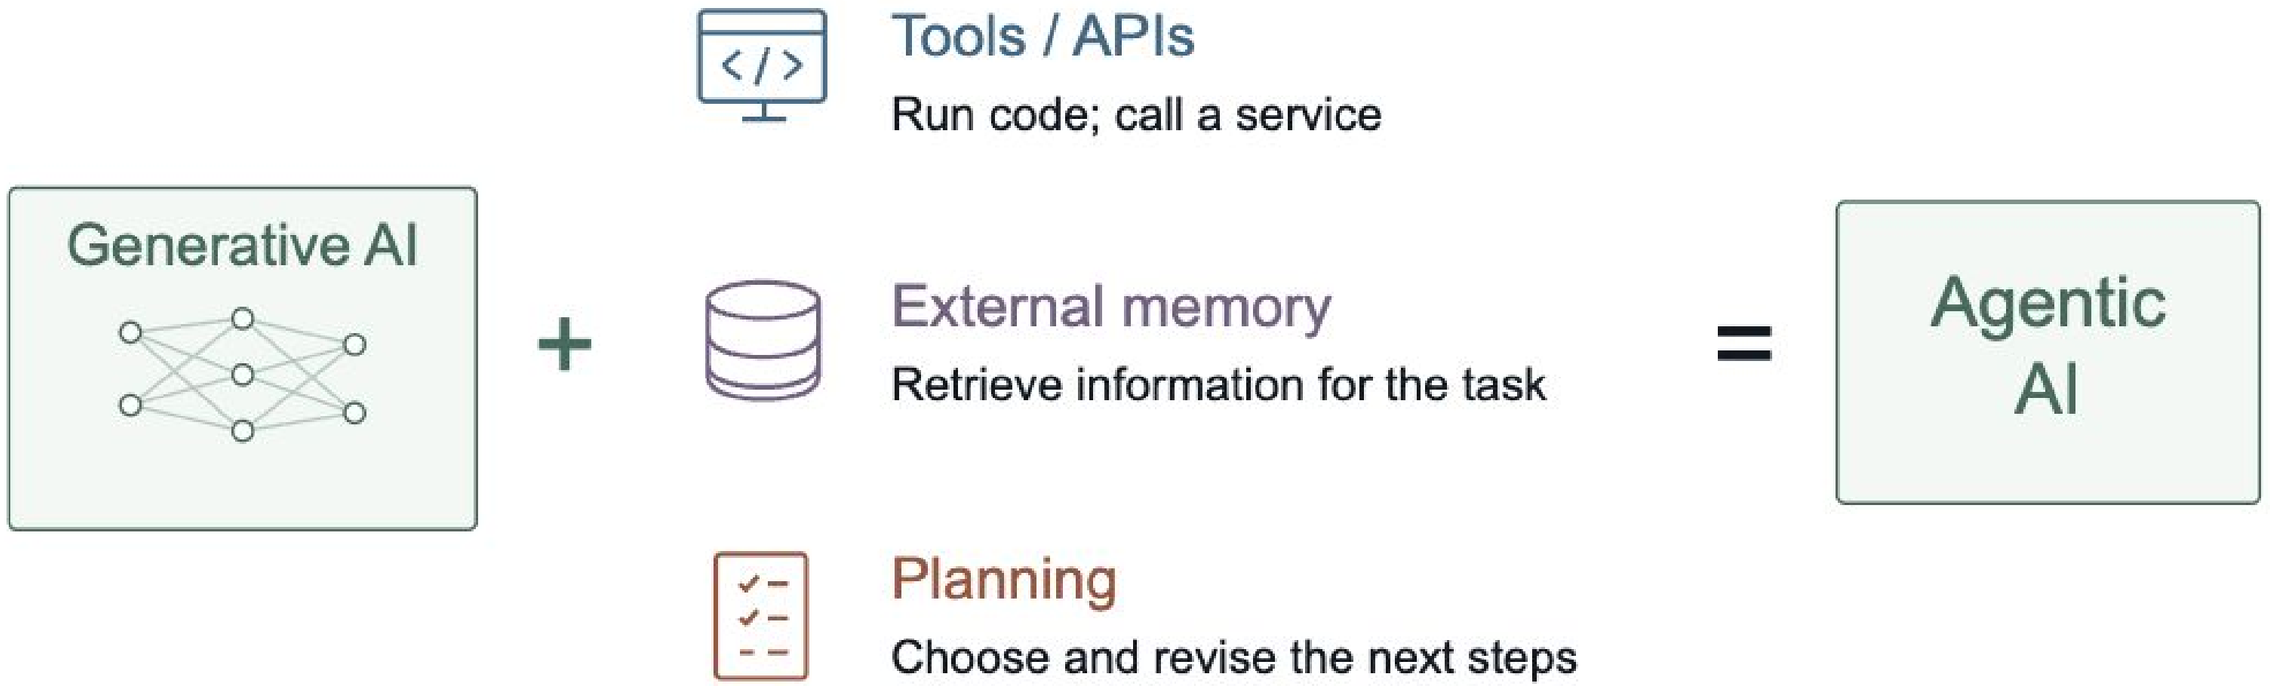

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What is an AI agent?

An agent uses an AI model to interact with its environment and work toward a user-defined objective.

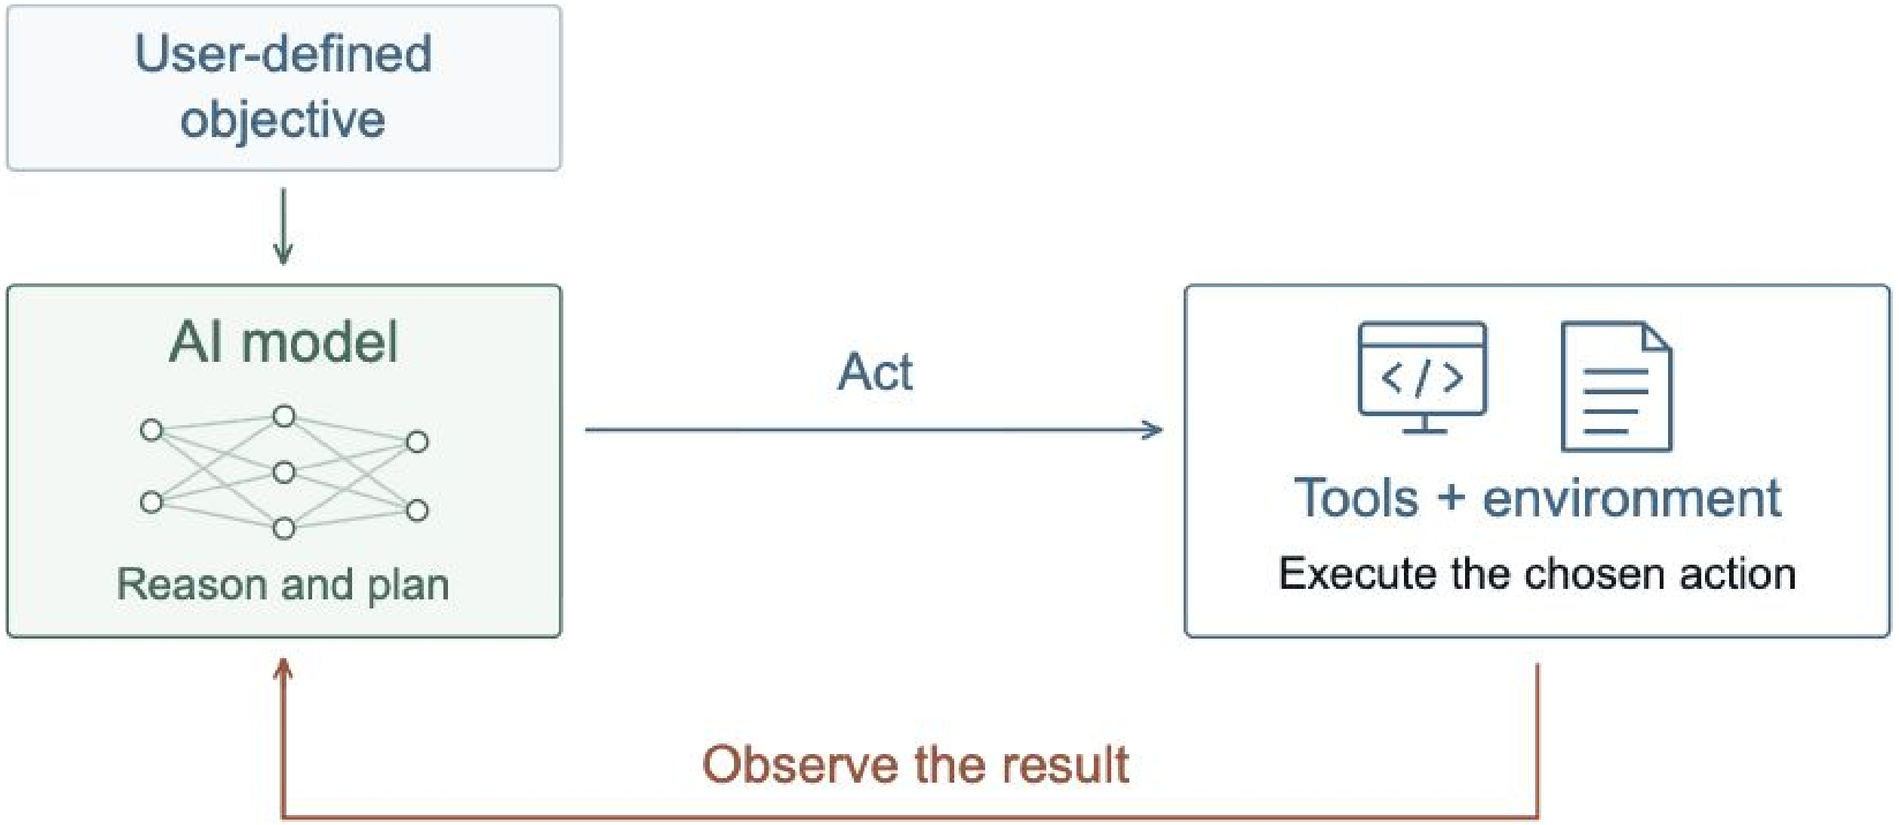

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

How are today's AI models built?

Let's look at the generative model inside the agentic application.

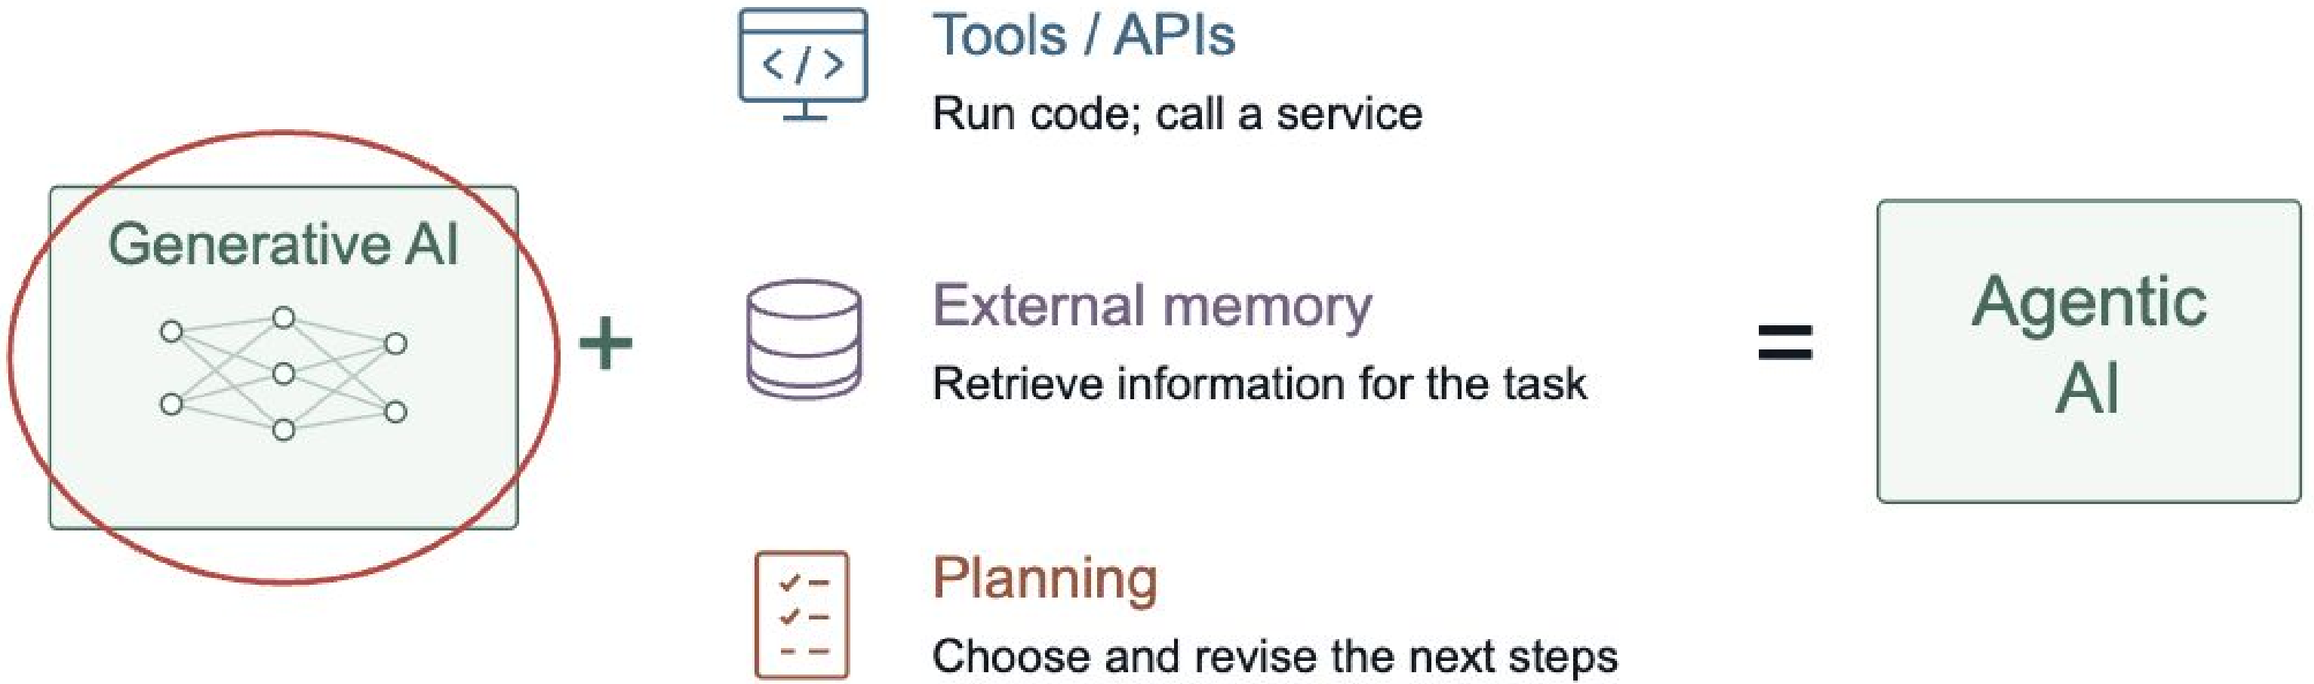

How does this model learn to produce useful outputs?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

From image classification to language models

The training idea stays the same: predict, compare, and adjust parameters.

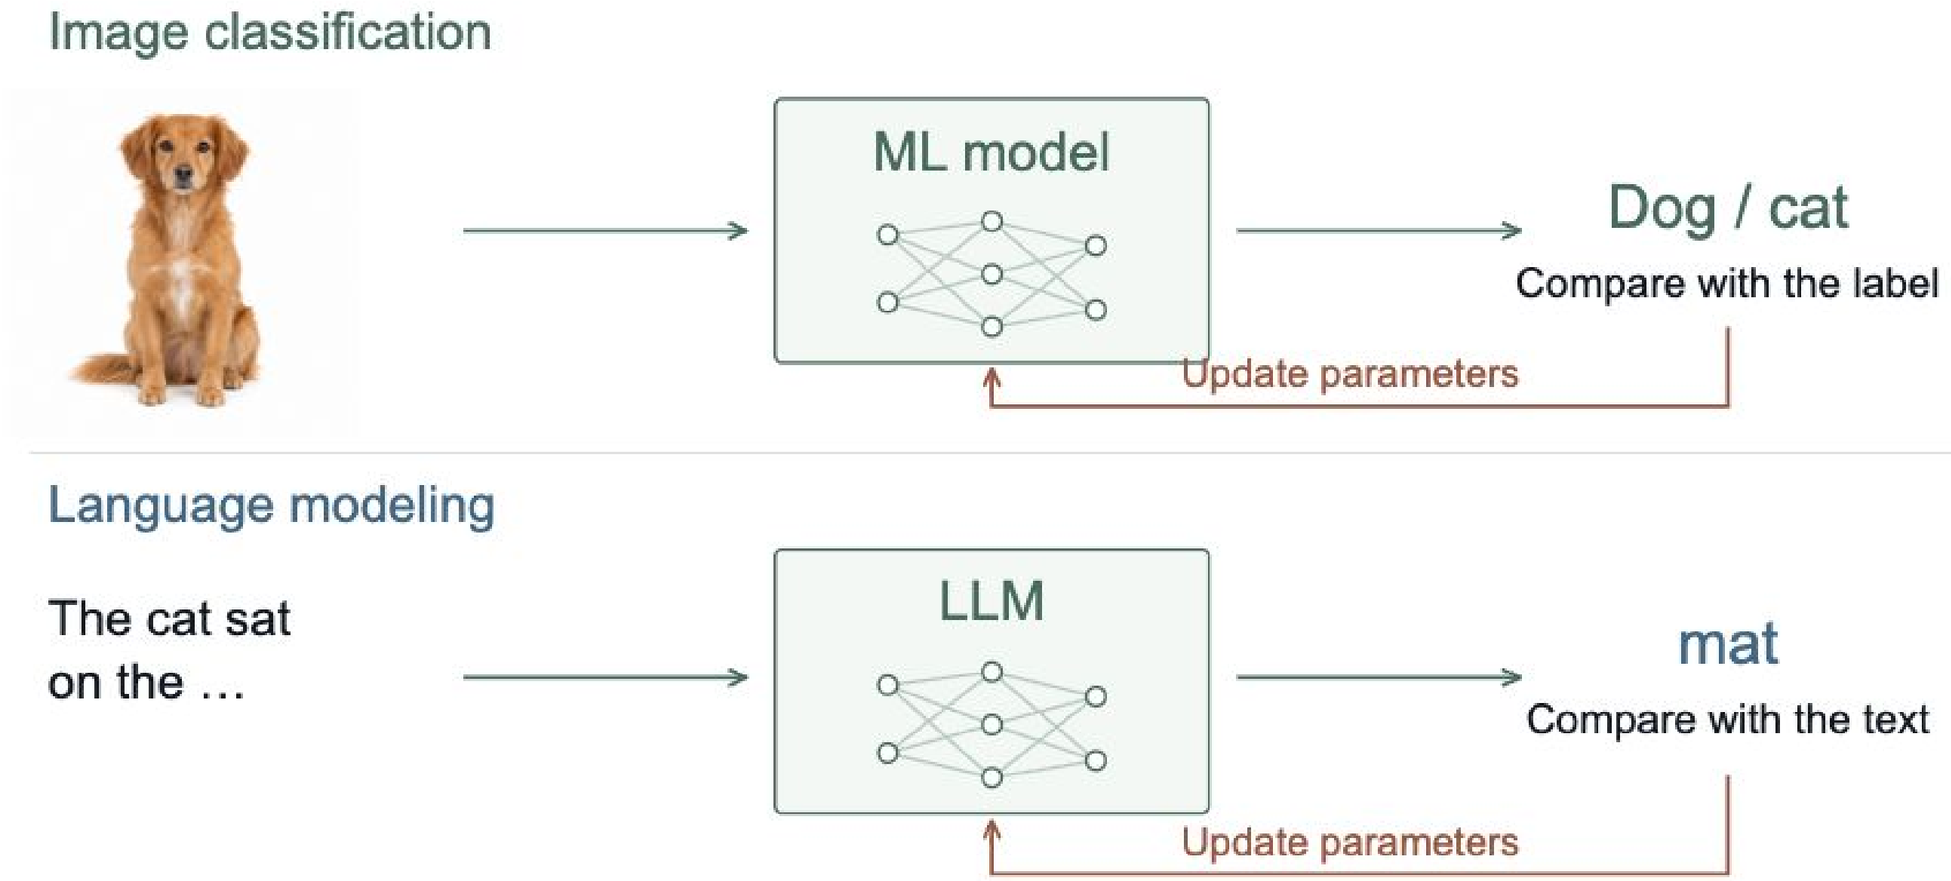

For a language model, what provides the learning signal?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Three stages of model training

A model can learn from text, labeled responses, and feedback.

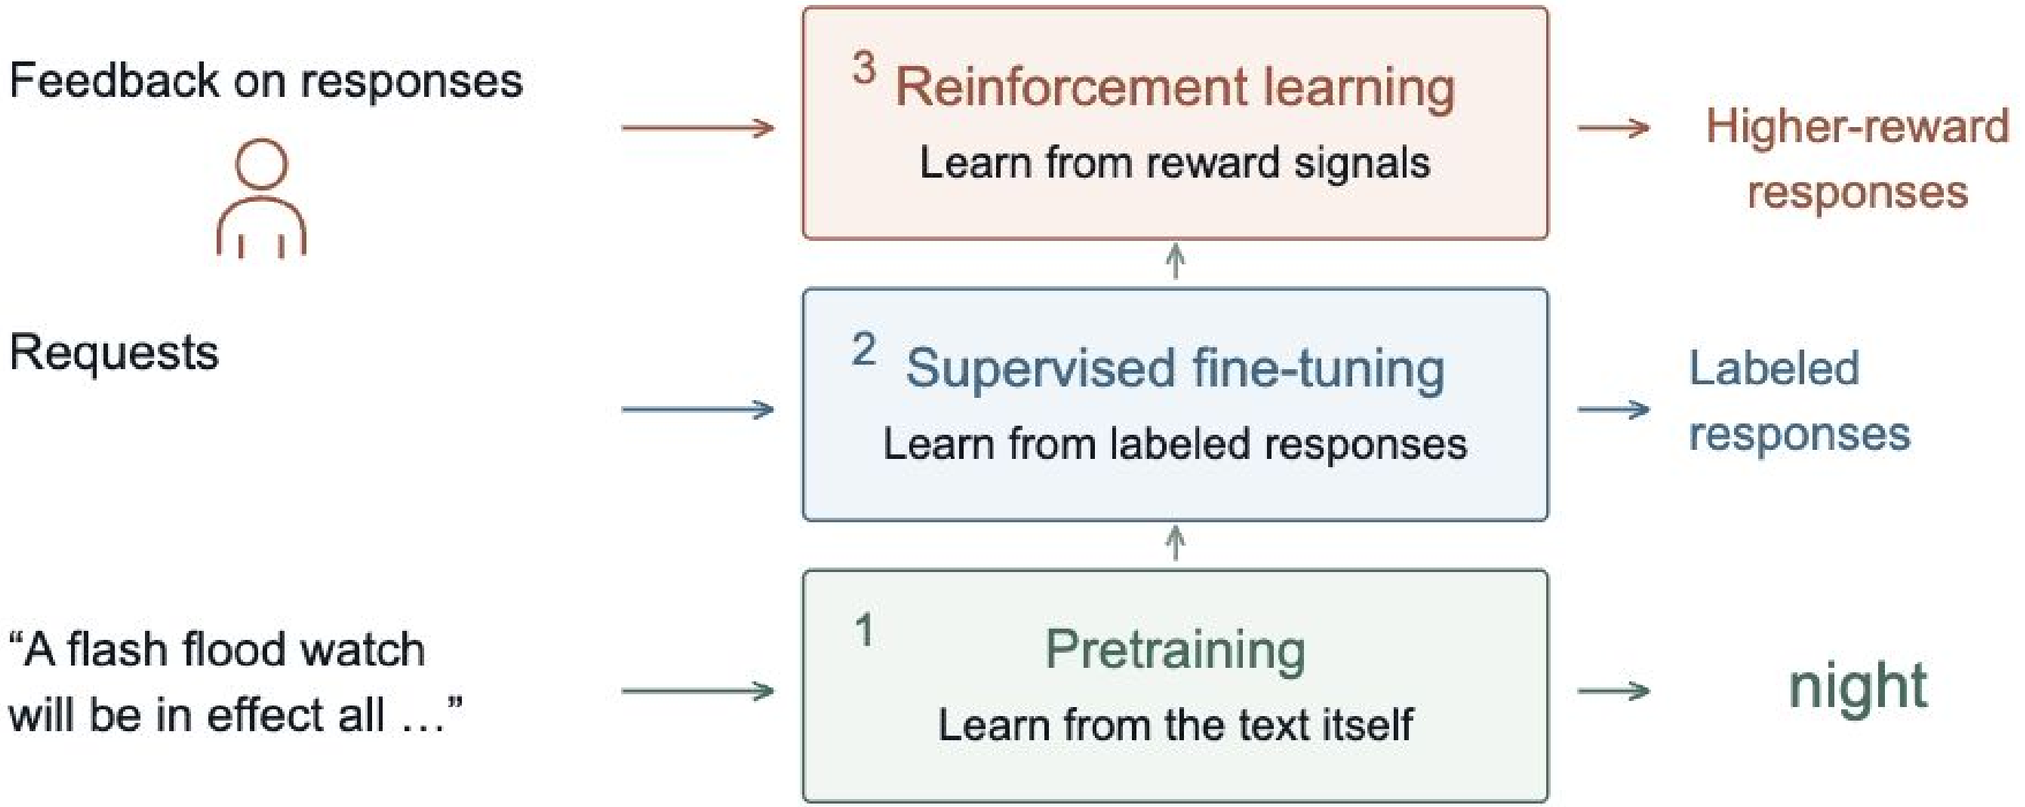

Start with pretraining, then shape how the model responds.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Phase 1: Pretraining

Learn patterns in large text collections by predicting what comes next.

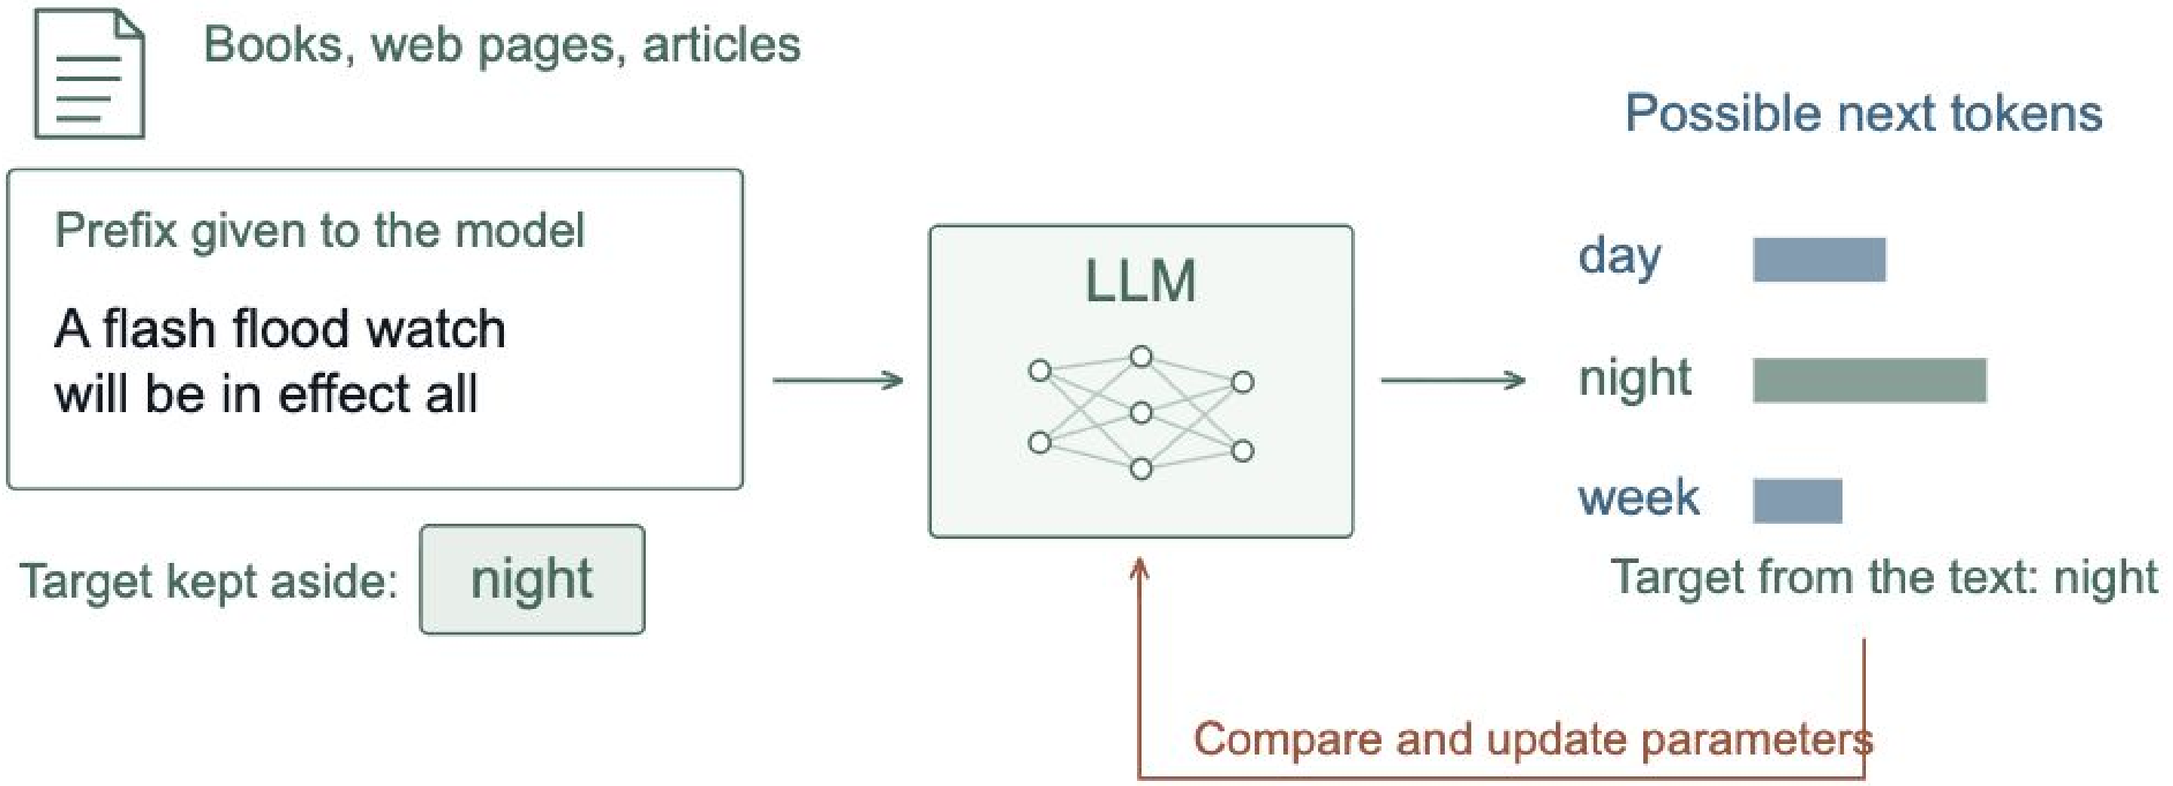

The original text supplies the answer. This is self-supervised learning.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Tokens

What does the model receive when we give it text?

DoiT is a great company to work for.

↓

Tokens	Do	i	T	is	a	great	company	to	work	for	.
Token IDs	528	72	51	318	257	1049	1664	284	670	329	13

Text → tokenizer → a sequence of numbers

A token can be a word, part of a word, or punctuation.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Why use tokens instead of whole words?

Reusable pieces balance vocabulary size and sequence length.

unbelievable → un | believ | able

Handle unfamiliar words

Build them from reusable pieces.

Keep sequences manageable

Use fewer pieces than individual characters.

Context window

Input tokens | Generated tokens

The model has room for a limited number of tokens in one call.

Now return to training: how do we teach the model to answer a request?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Phase 2: Supervised fine-tuning

Train on requests paired with labeled responses.

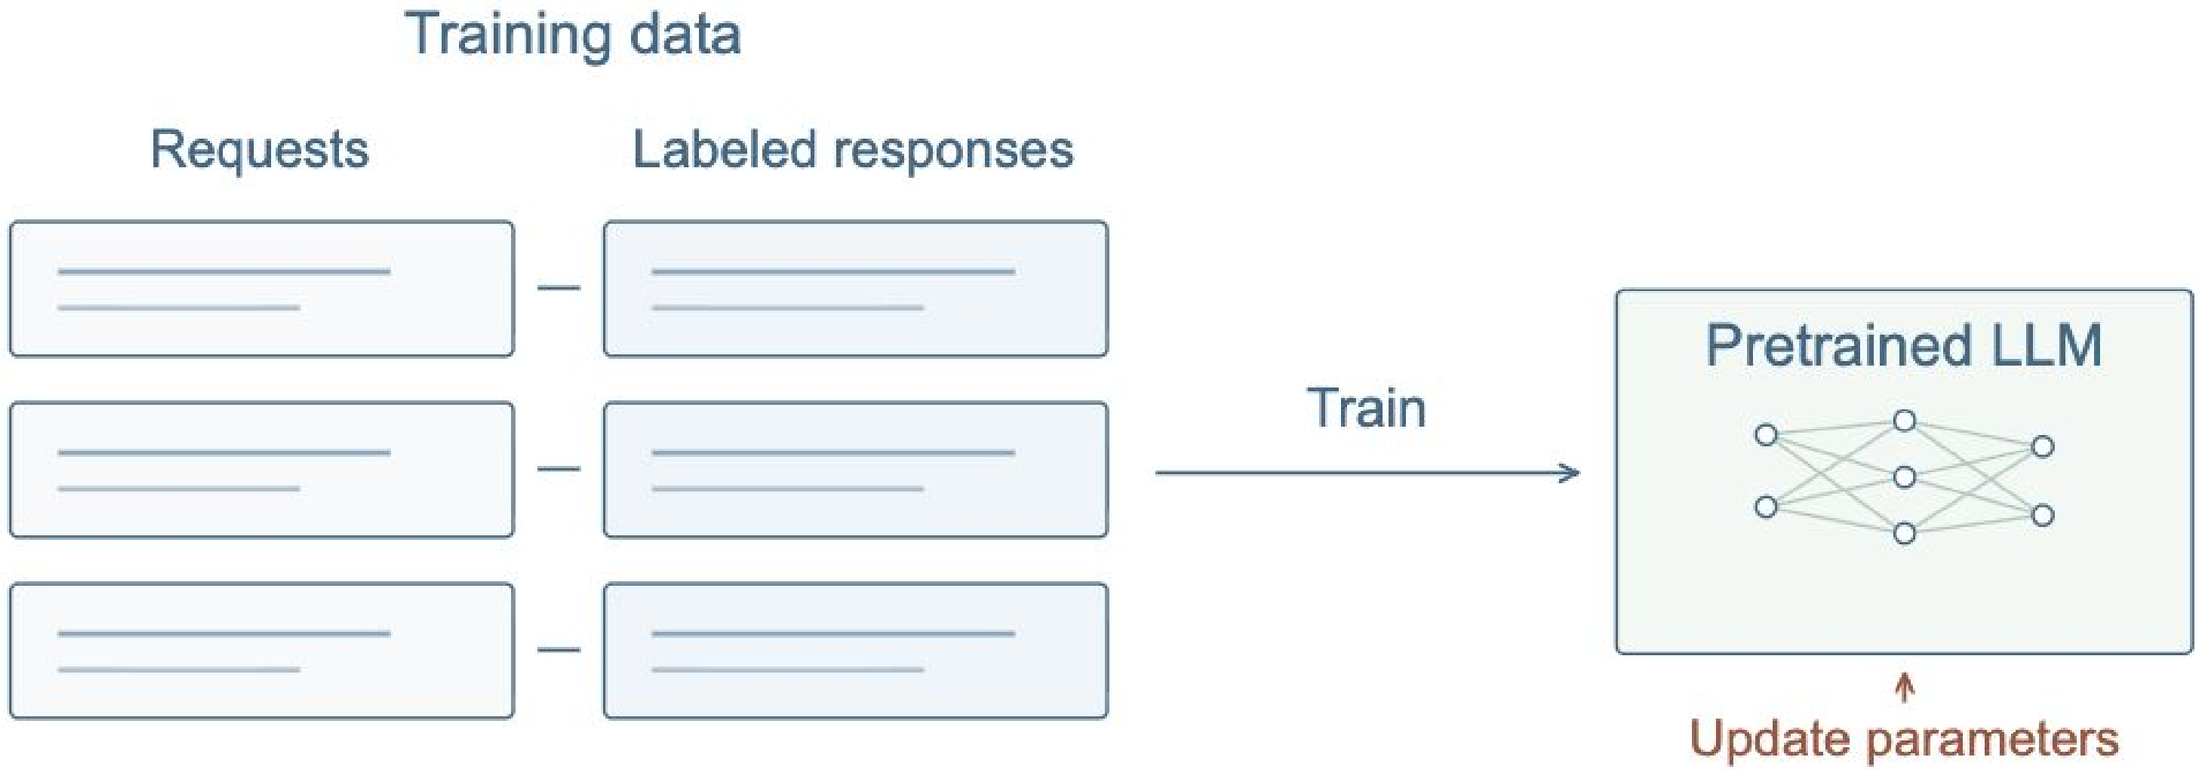

Each labeled response gives the model a target to learn from.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Phase 3: Reinforcement learning

Generate responses, score their quality, and use the reward to update the model.

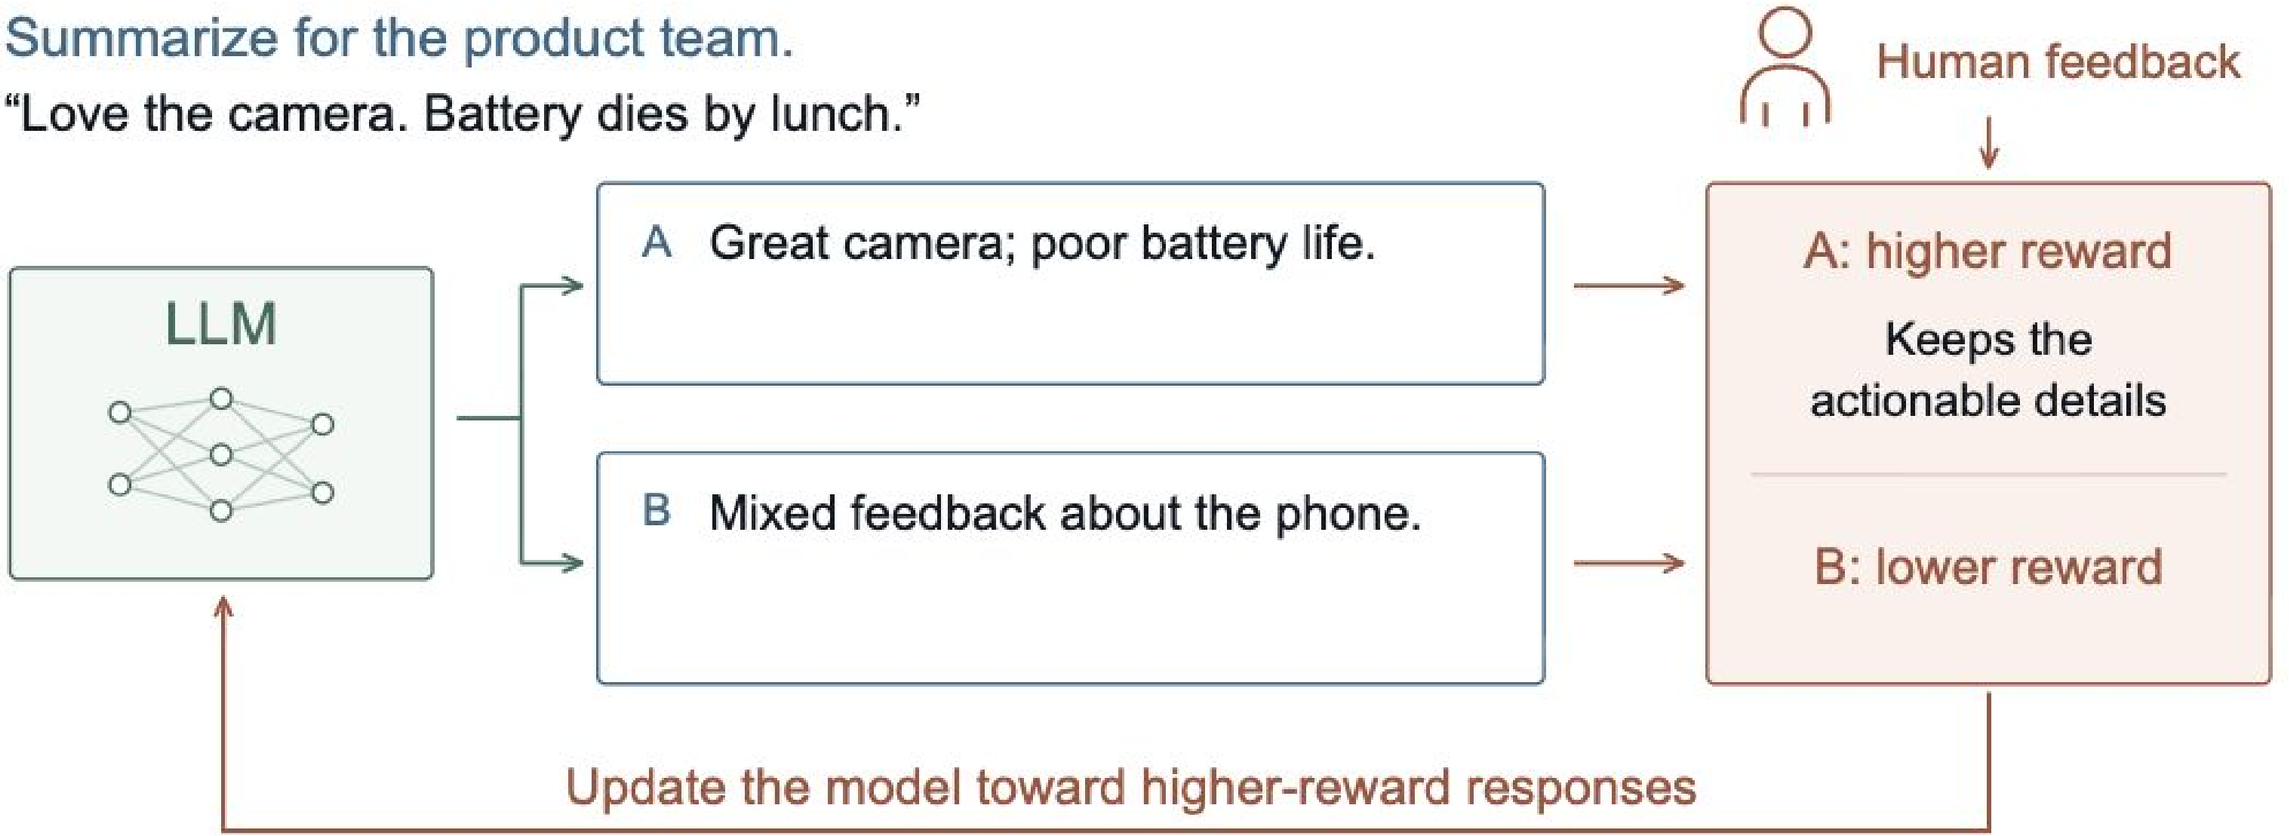

Here, human feedback tells us which response is more useful.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Phase 3: RLHF

Reinforcement learning from human feedback turns preferences into a training signal.

1. People teach the reward model

2. Its scores guide the LLM's training

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Post-training: Phases 2 & 3

After pretraining, labeled responses and feedback shape the model's behavior.

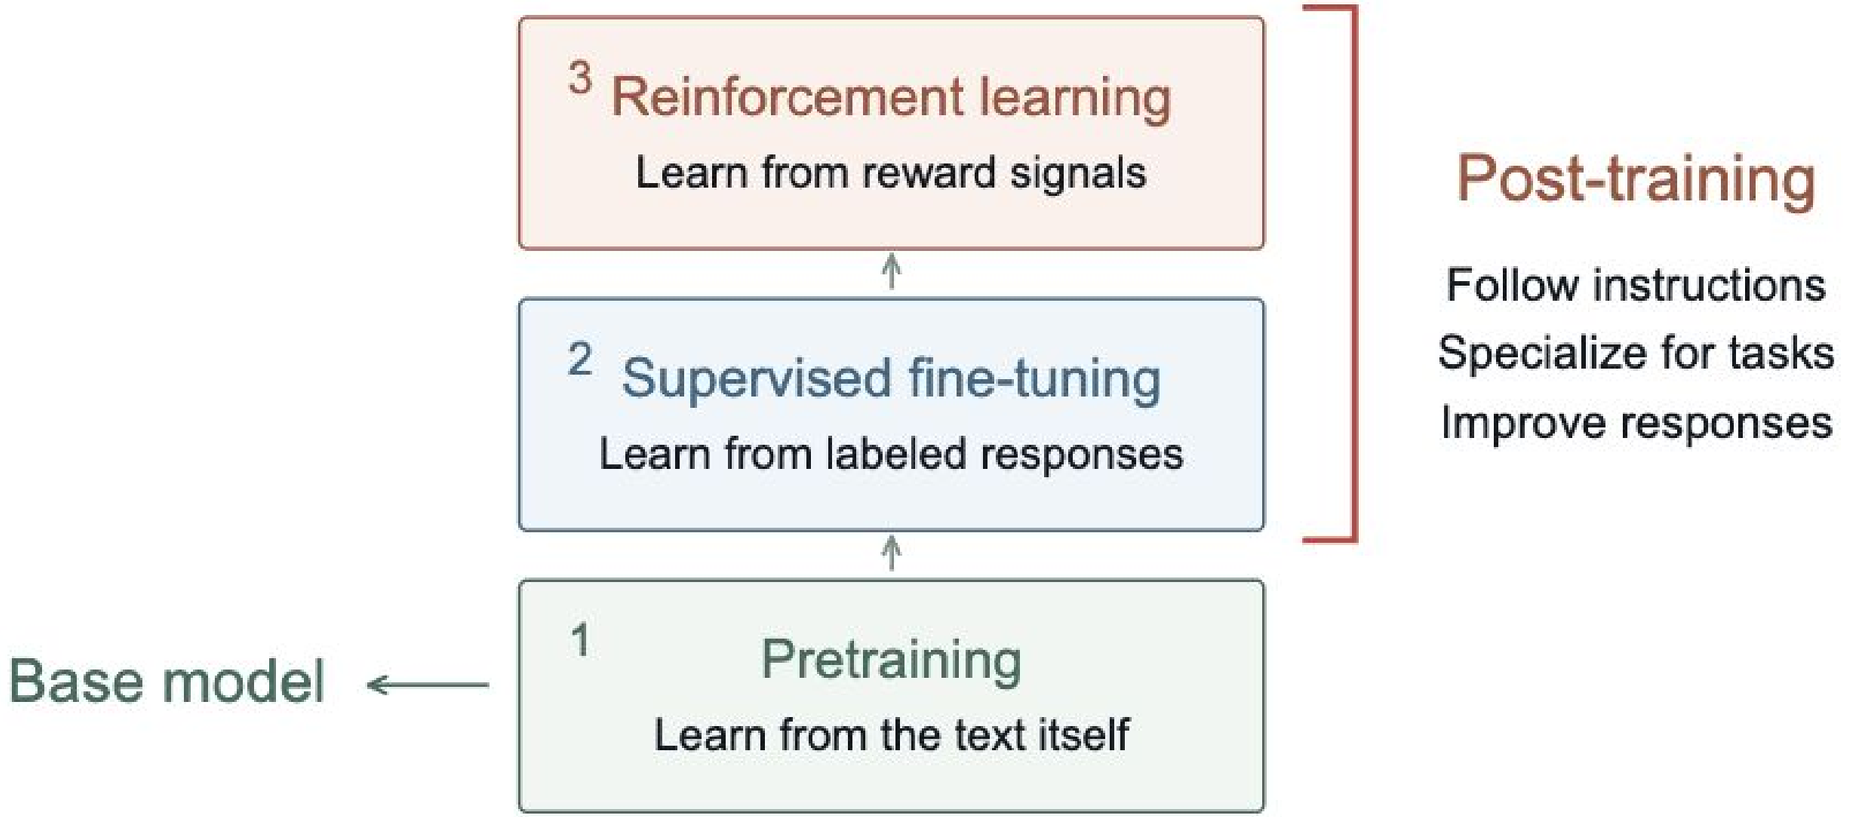

The model keeps learning, but the training signals and objectives change.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Why do models have a context limit?

The context window limits how much information a model can use in one call.

Short input

Longer input

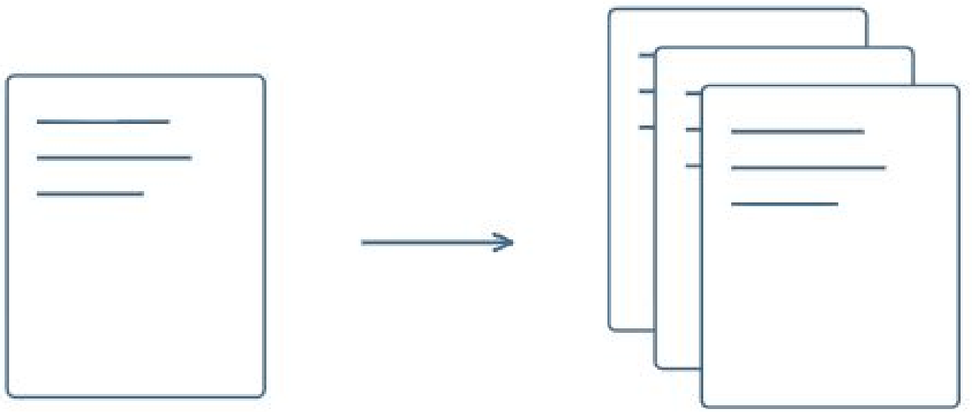

More information to process

More resources needed

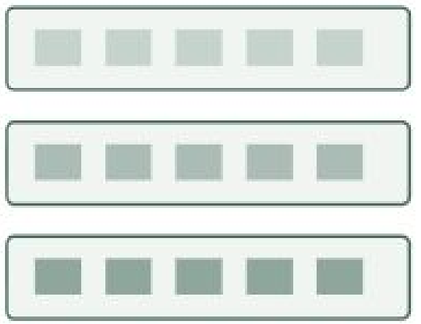

Computing power and memory

Longer inputs take more resources to process.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What happens after training?

We use the trained model to generate outputs for new inputs.

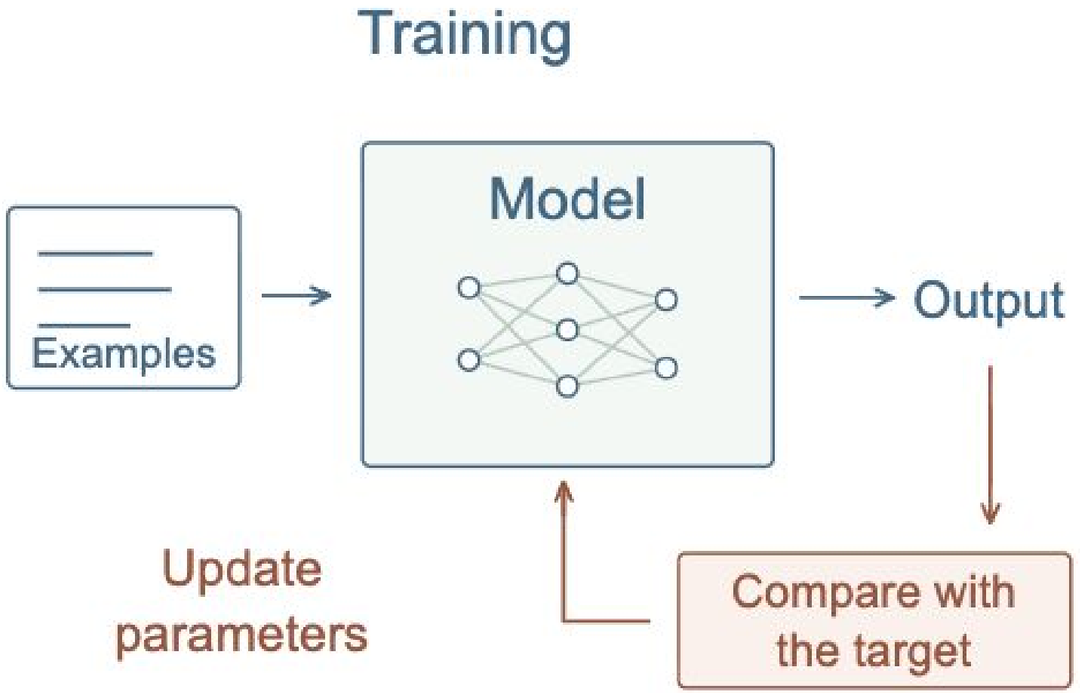

Examples and feedback change weights.

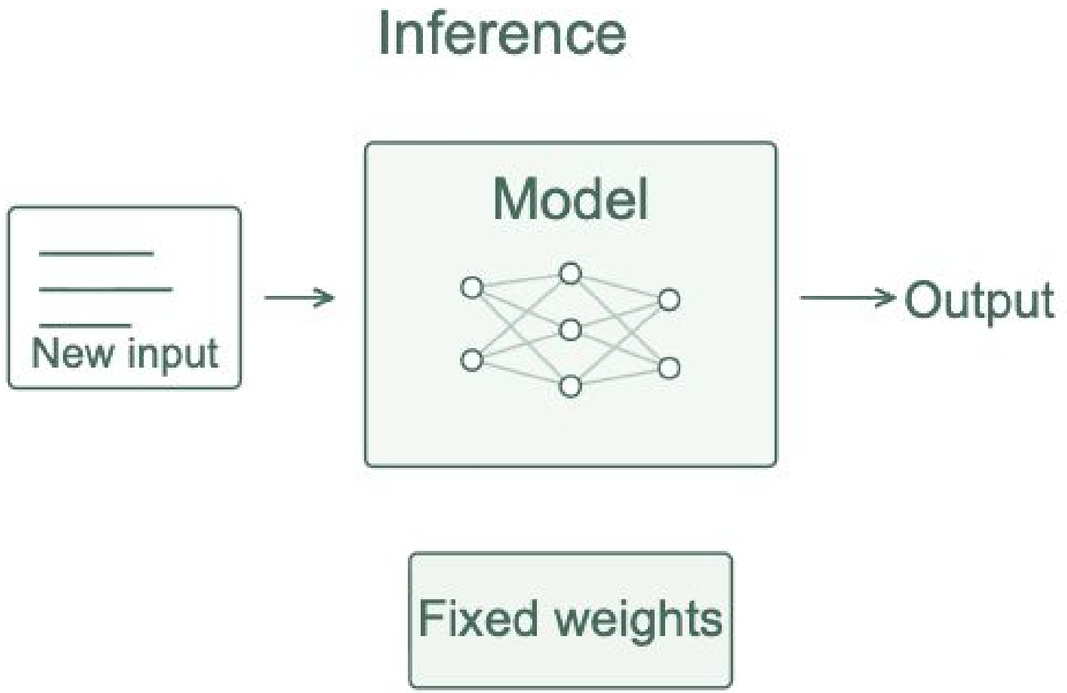

New inputs use the learned weights.

Now we can ask: how good are the outputs?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

How do we evaluate a model?

There is no single “#1 model” across every task and test setup.

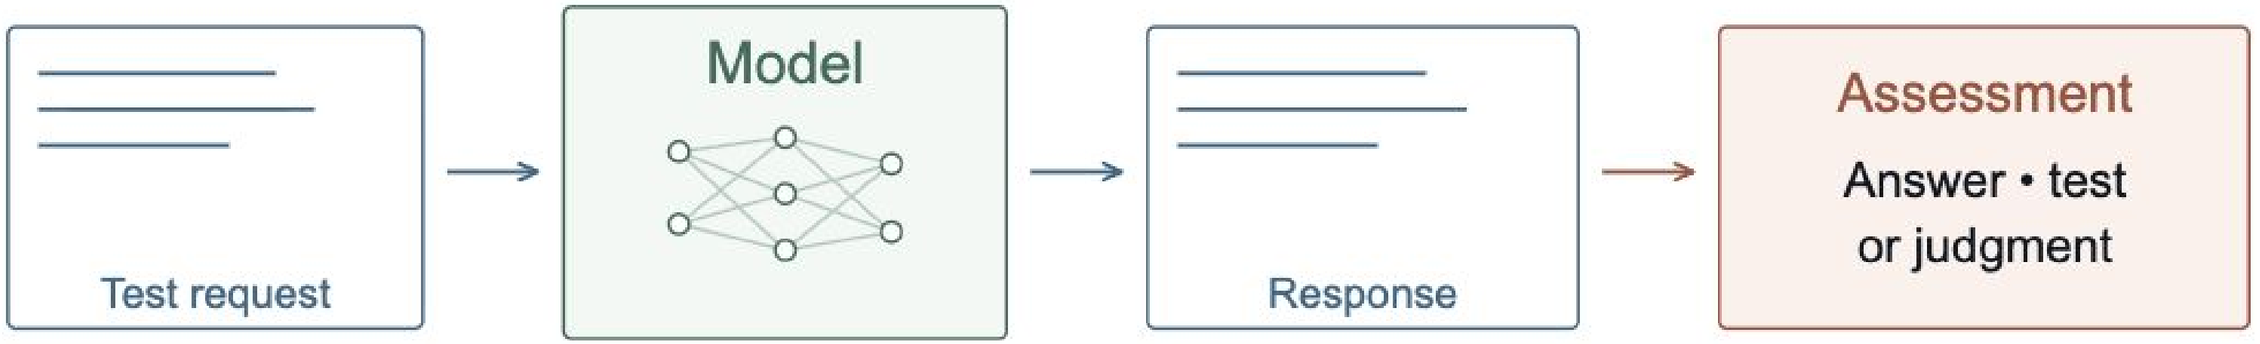

A task from the test set

Which task?

The work being tested

Which data?

The examples selected

Which setup?

Prompts, tools, budget

Evaluation measures performance on a defined test.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Evaluation during training

Different tests reveal different capabilities as the model develops.

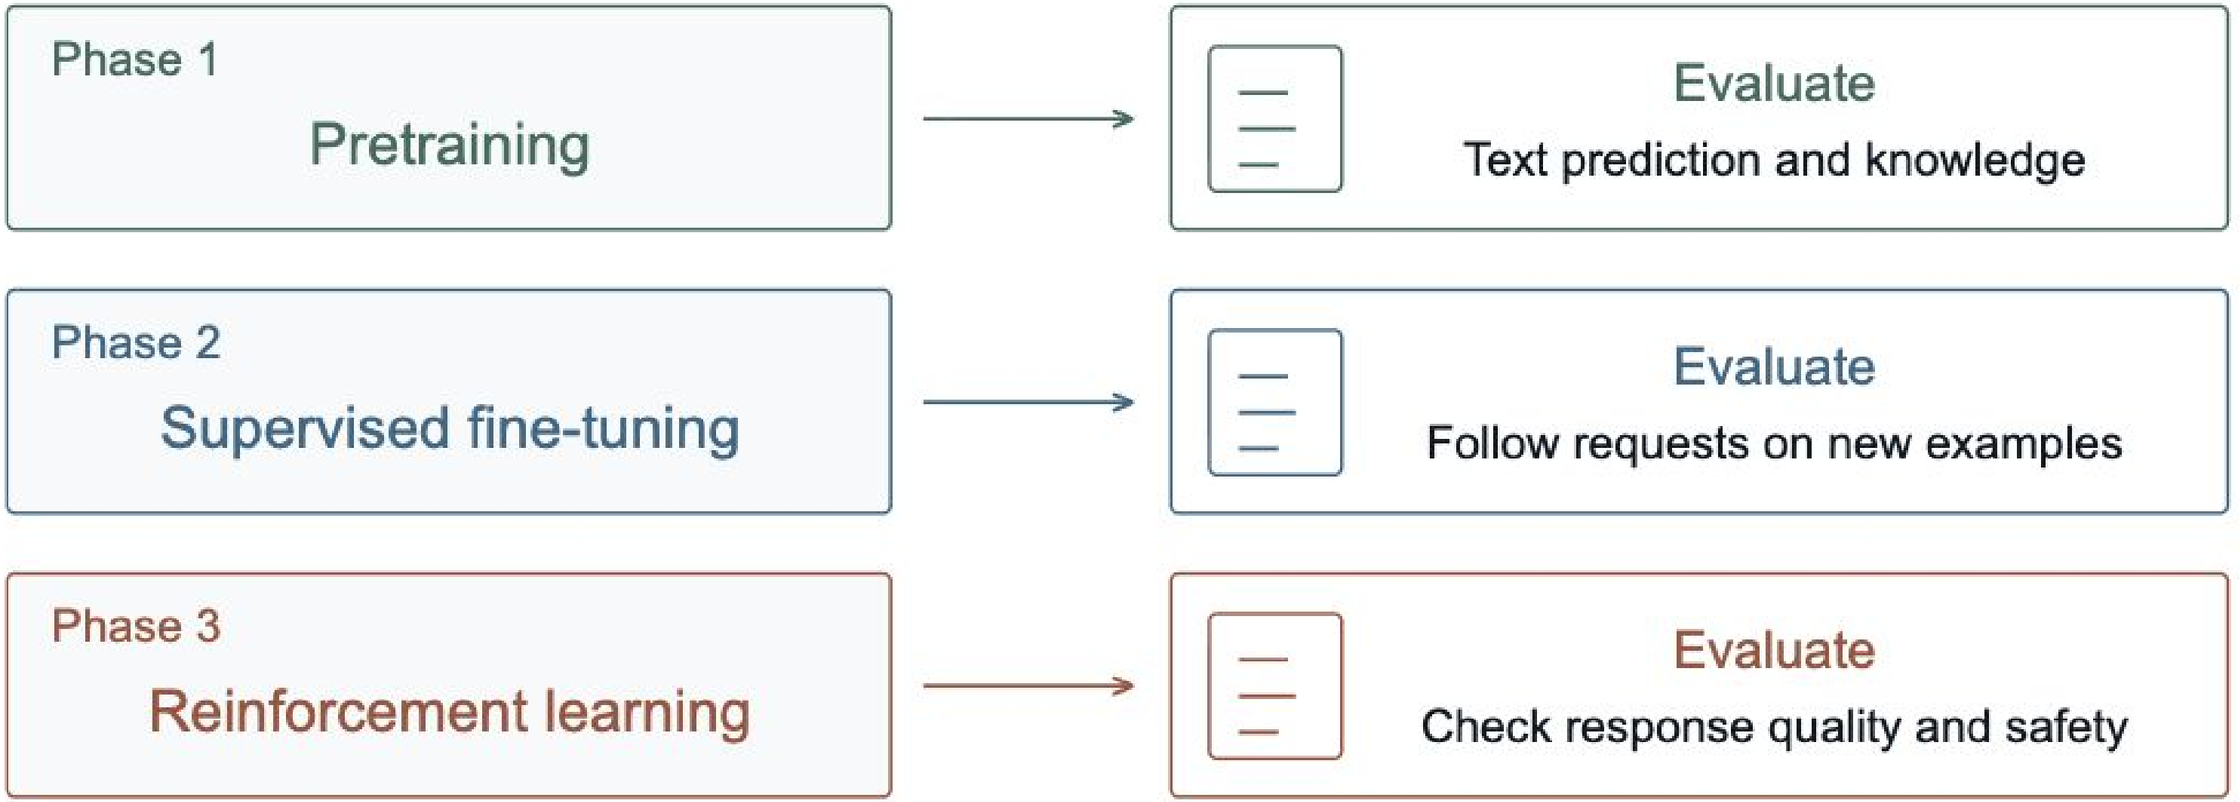

Tests measure progress throughout training.

A final model is evaluated across multiple kinds of work.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What is a benchmark?

A shared set of tasks and scoring rules used to compare models.

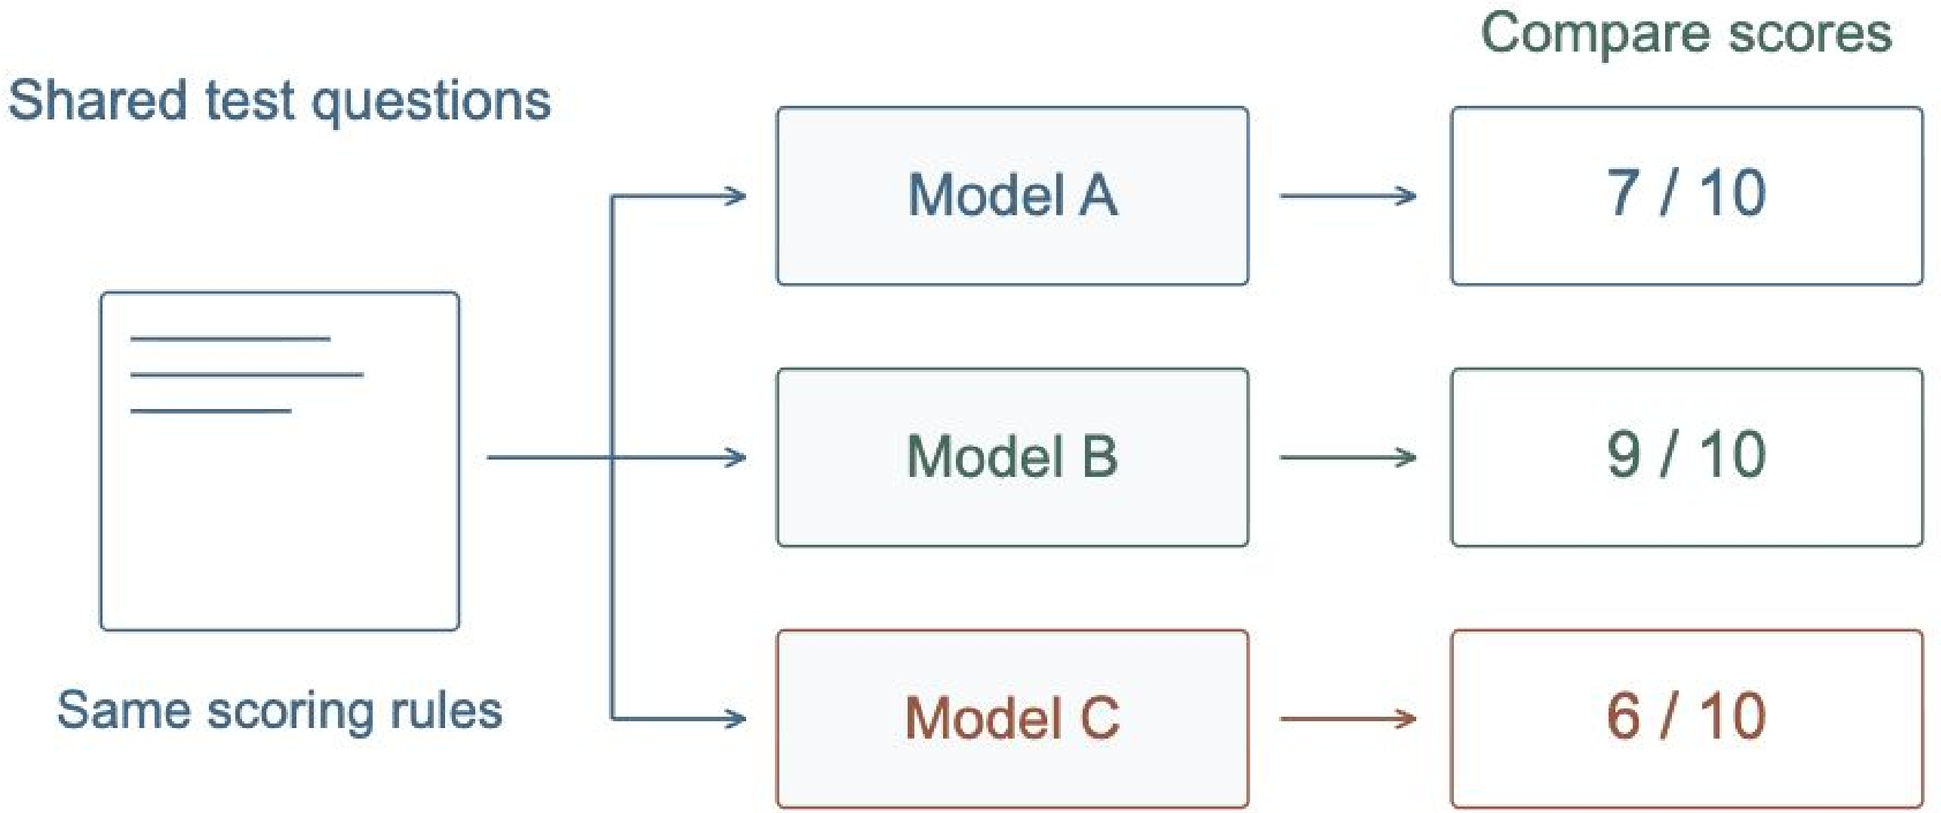

For example: compare how many of 10 questions each model answers correctly.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Two ways to classify a benchmark

The question source and scoring method are separate choices.

Where do the questions come from?

Static questions

Reuse the same test questions.

Live / refreshed

Bring in new test questions.

How are responses scored?

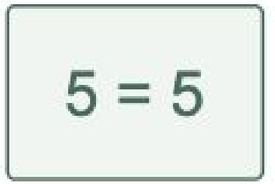

Ground truth

Check an answer or run a test.

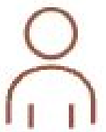

Human preference

People judge which response is better.

Combine the two choices to get four benchmark types.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Four types of benchmarks

The same task can be evaluated with different question and scoring choices.

	Static questions	Live / refreshed
Ground truth Checkable result	**MMLU-Pro** Fixed questions answer keys	**LiveBench** New questions checkable answers
Human preference People compare	**A fixed prompt set** Reviewers compare responses	**Arena Battle Mode** New prompts user votes

Read both the test questions and the scoring method before comparing scores.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What do benchmarks measure?

Choose tests that resemble the work you care about.

Knowledge & reasoning

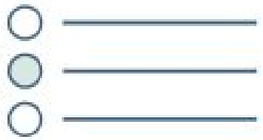

Answer across subjects

MMLU-Pro

Coding

Fix a repository issue

SWE-bench

Math & science

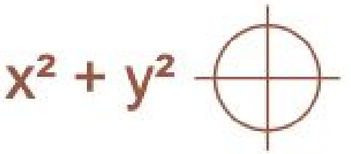

Solve and explain problems

LiveBench math / GPQA

Professional work

Produce a work deliverable

GDPval

Safety & adversarial

Handle harmful requests

HarmBench

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Human-preference evaluation: Arena

In Battle Mode, users compare responses before the model names are revealed.

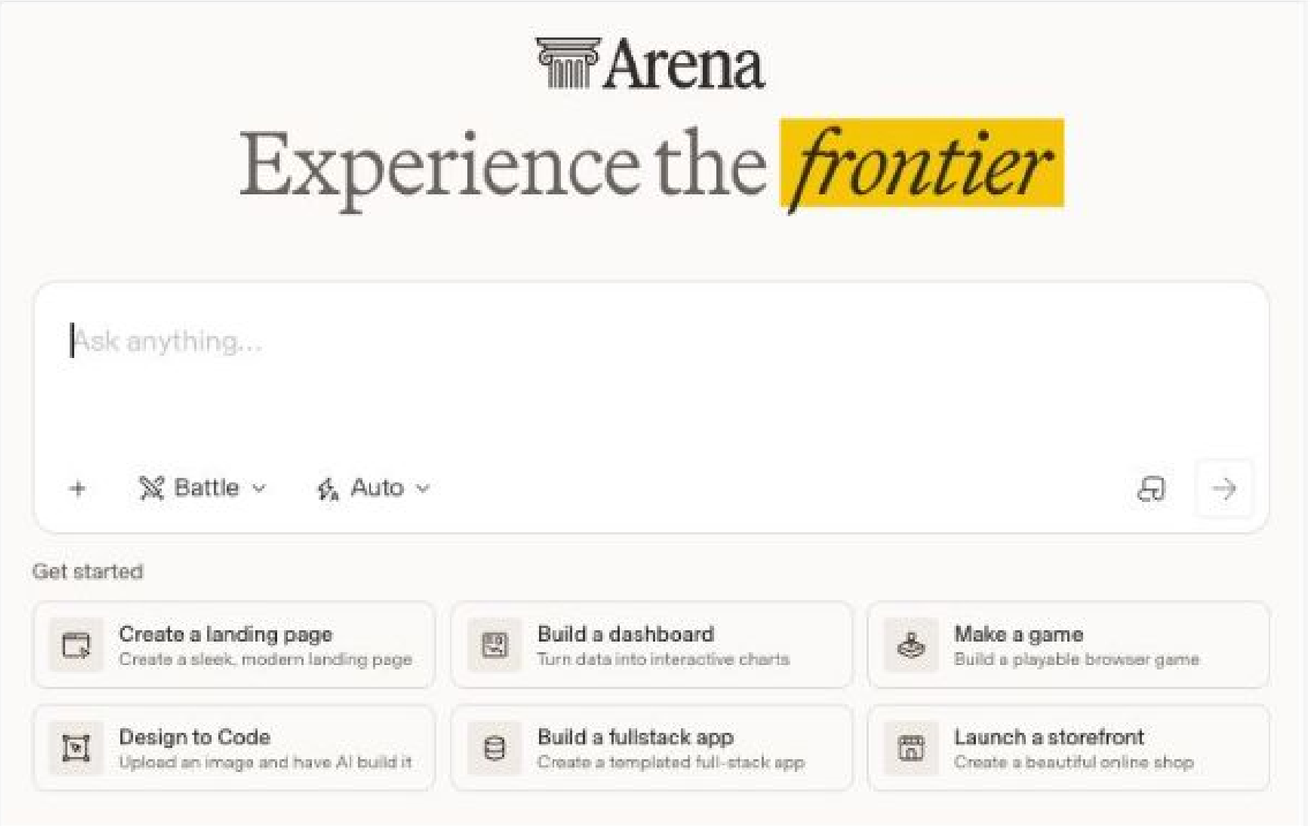

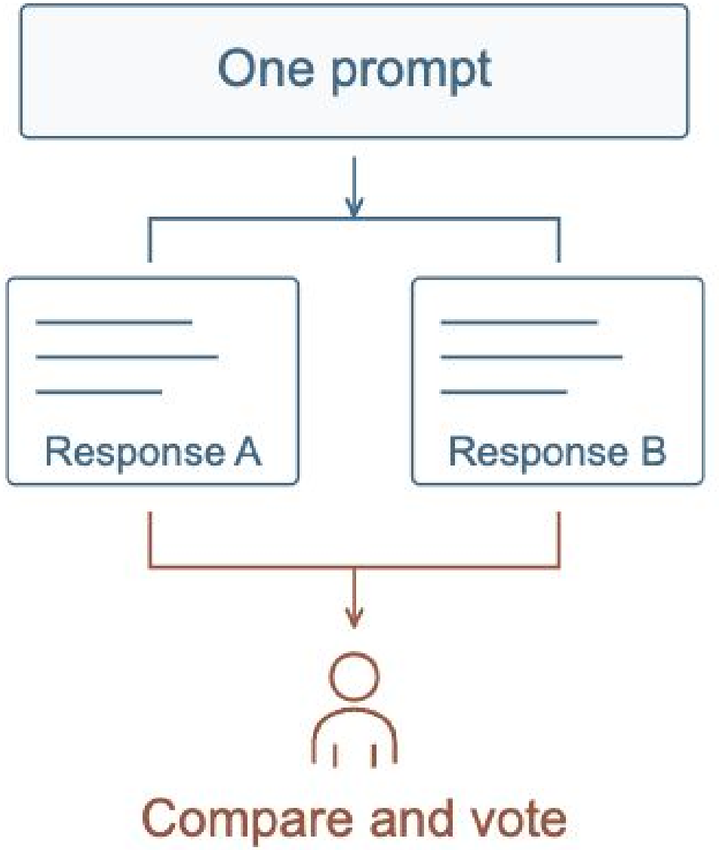

The result measures which responses people prefer.

Screenshot: Arena (October 2026)

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

LevelUp Lens™

All AI benchmarks are biased, but some are more useful than others.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Model evaluation and product evaluation

Is the model response right? Does the whole application work?

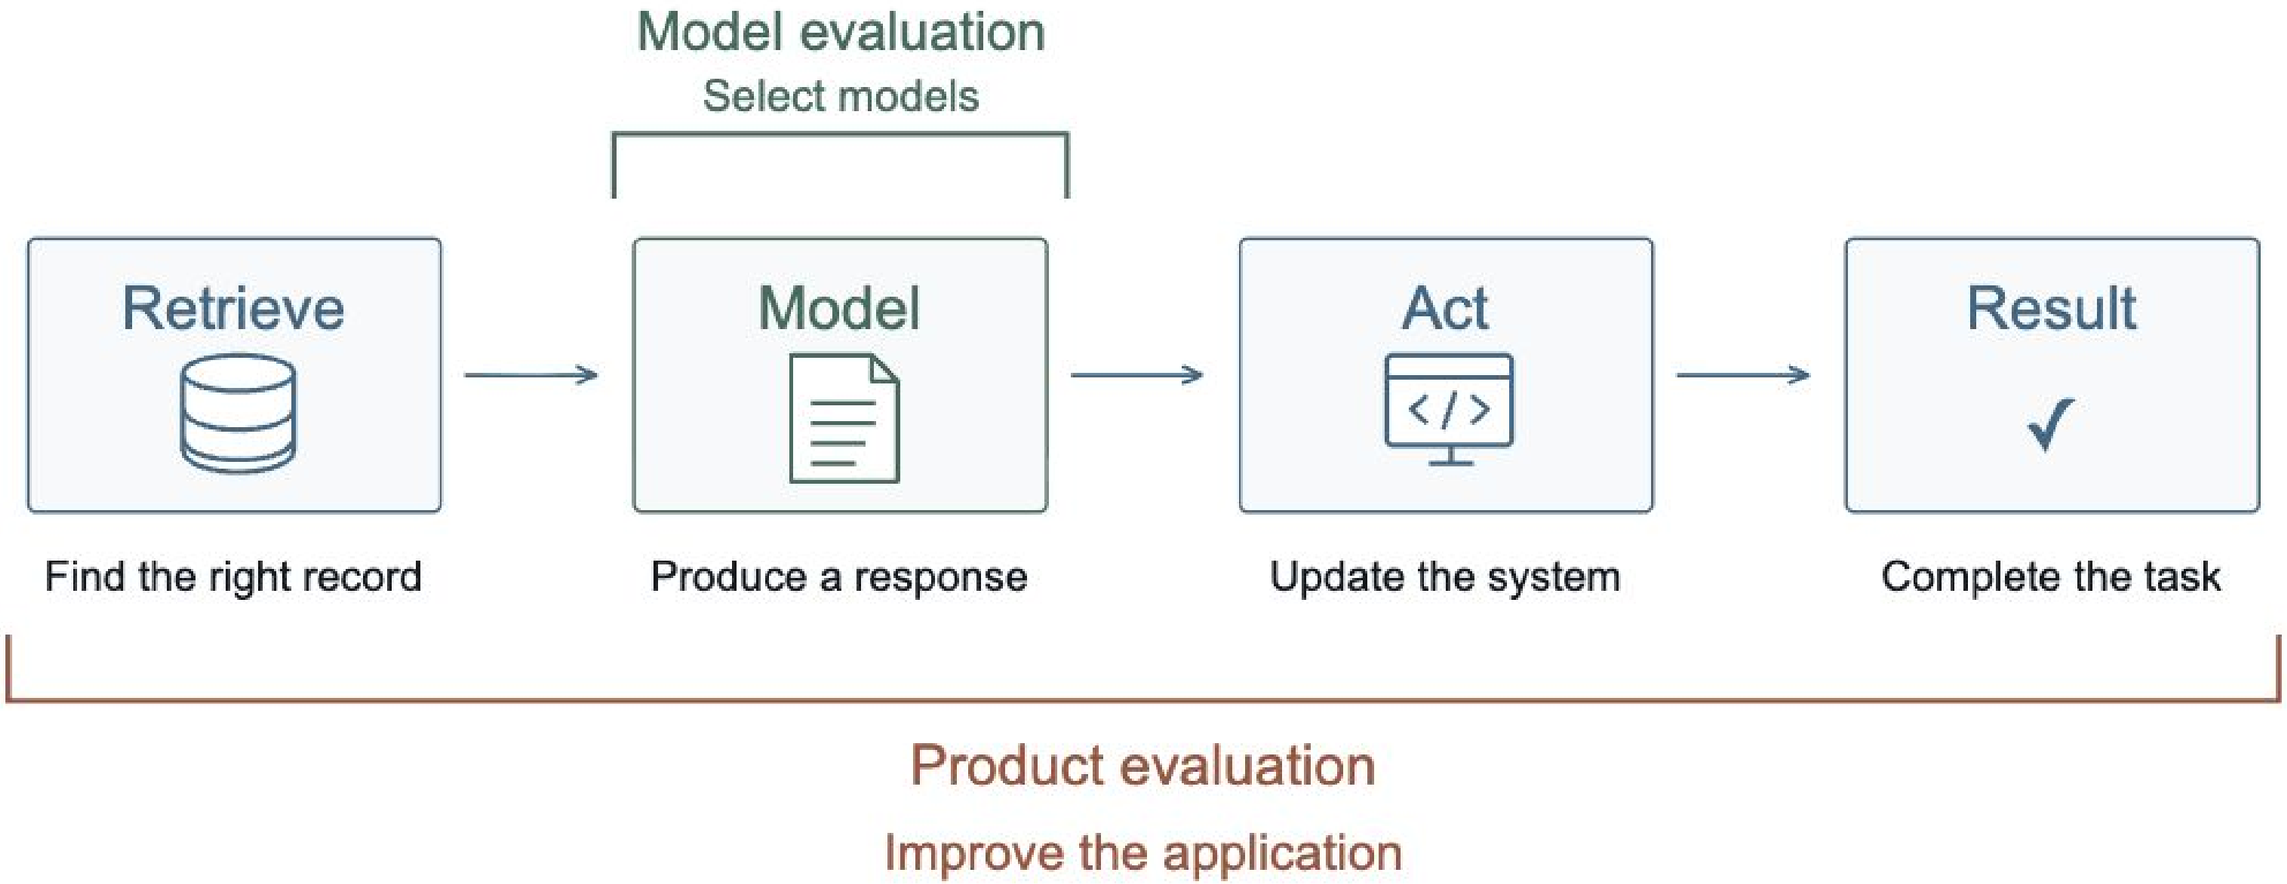

Most enterprise builders improve how the model is applied.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

LevelUp Lens™

Model evals

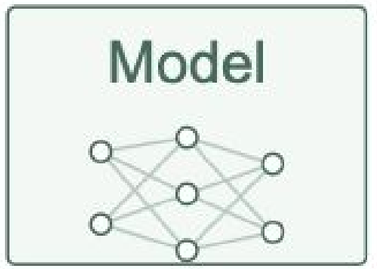

Test performance

Product evals

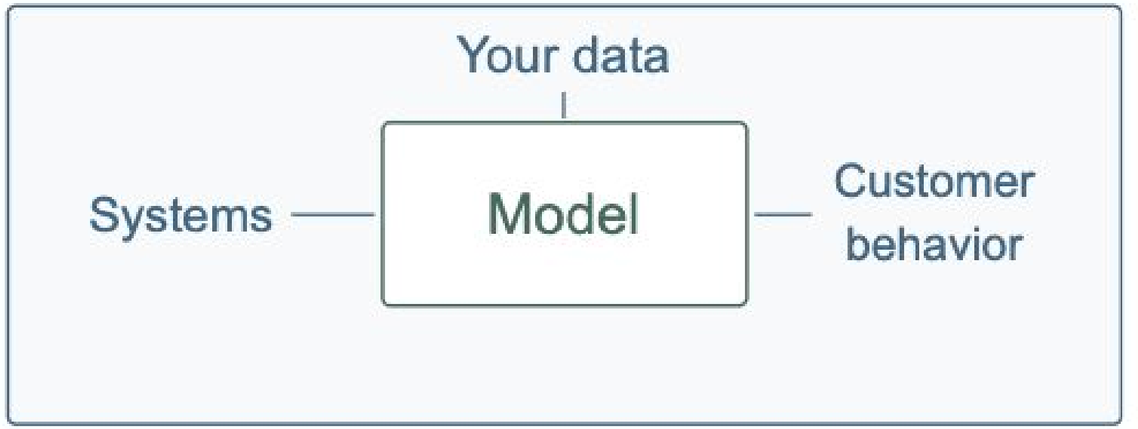

Model evals are only a starting point: product evals show whether AI works with your data, your systems, and the way your customers use it.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

How do we make sense of model advances?

We know how to evaluate a model.

When a new model arrives, what has improved?

The Input-Output Framework™

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The Input-Output Framework™

Start with the model we already know.

Image input

Learned model

Image classifier

Uses learned patterns

Category output

✓ Dog

Cat

Chooses one label

A category label is one output. What else can a model produce?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The Input-Output Framework™

Different models produce different kinds of output.

Input	Model	Output	Examples
Text	Language model	Text	GPT-3.5 Turbo
Text	Video model	▶	Veo 2
Text	Audio model	〰	ElevenLabs Multilingual v2
Text	Image model		Nano Banana

Logos: OpenAI · Google · ElevenLabs

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Reasoning models

Same text input. More reasoning before the answer.

OpenAI o3

DeepSeek-R1

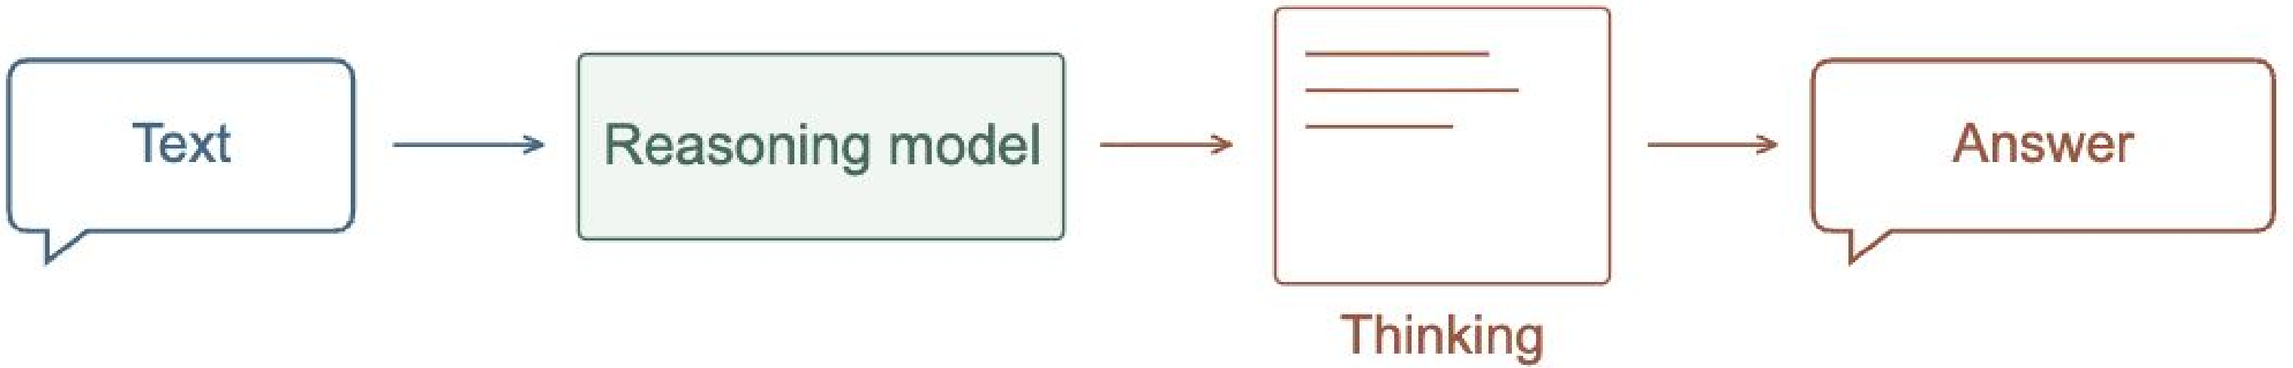

More computation before the final answer.

Logos: OpenAI · DeepSeek

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Reasoning mode and effort

The effort setting guides how much work the model puts into a response.

You change the effort for this response, not the model's training.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Reasoning models: a simple example

Same question. Same answer. Extra steps in between.

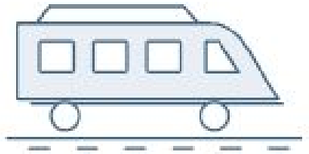

A train leaves at 3:00 PM and arrives at 5:30 PM.
How long is the trip?

Regular response

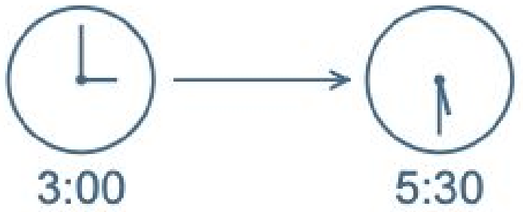

2 hours 30 minutes

Reasoning before answering

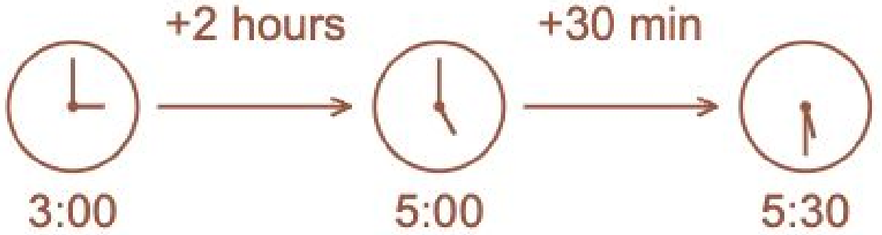

2 hours 30 minutes

Reasoning adds computation before the answer.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Reasoning, cost, and latency

More computation can help on harder problems.

Less computation

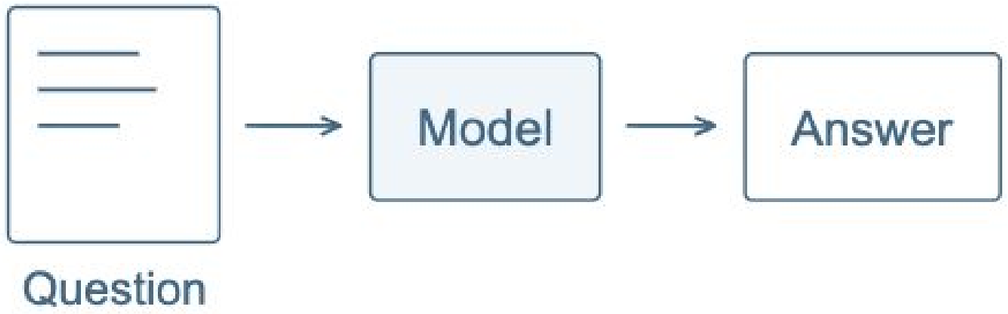

Often less time and cost

More computation

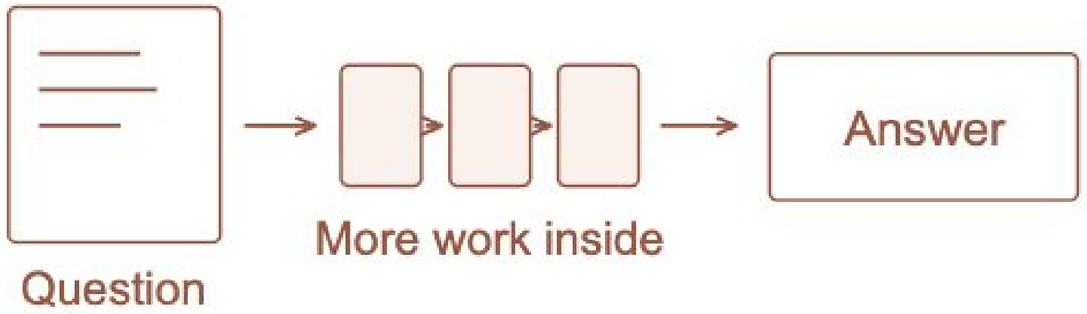

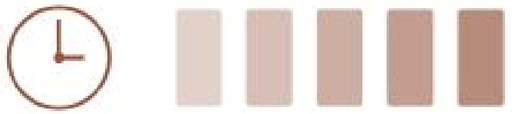

Can take more time and cost

Use extra reasoning where the improvement is worth the time and cost.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Decision models

Same text input. A choice from defined classes.

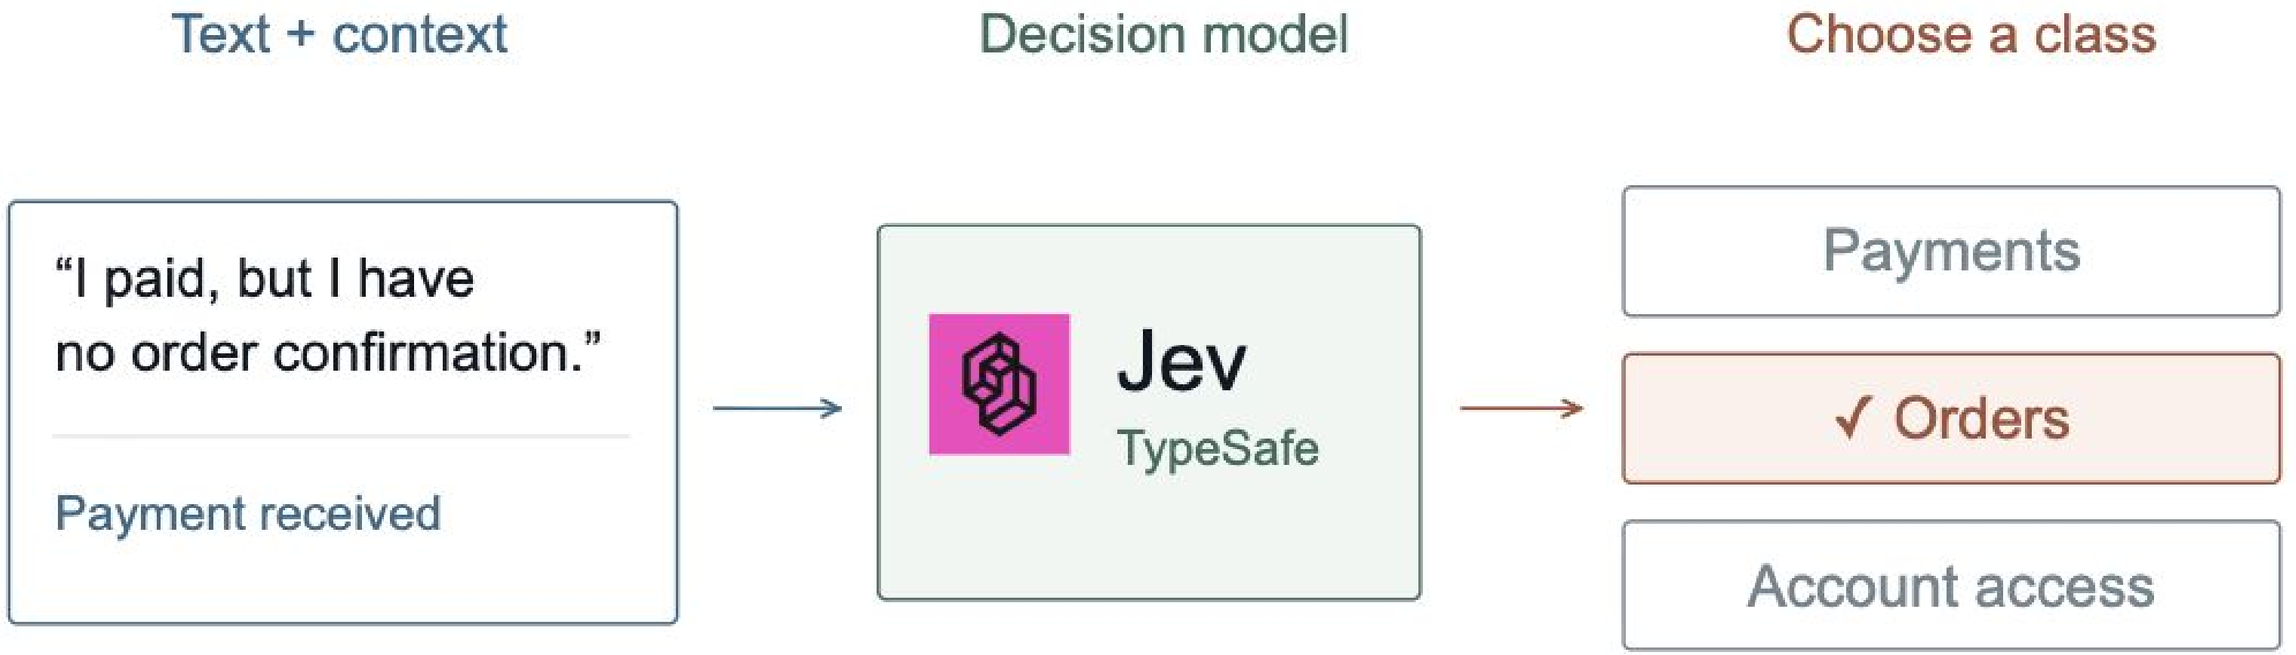

Your application sends the ticket to the Orders queue.

Logos: TypeSafe

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The Input-Output Framework™

Three ways to understand model improvements.

New capabilities

Text prompt →

New input-output combinations

Better task performance

→ Cat ×

Dog ✓

Do a particular task better

Lower cost or latency

Same useful result

Less time / cost

Make a known capability cheaper or faster

Science → Utility

More news-worthy

More use-worthy

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The Holy-Grail!

GPT-6 Astra

Opus 5.5

- Text
- Videos / Images
- Actions (Web tasks)

Logos: OpenAI · Anthropic

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

03 Choosing the right problems for AI

We understand what models can do. Which business problem should we tackle first?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What kinds of problems fit AI?

Look for work where writing every rule is difficult, but judging the result is possible.

- Hard to write every rule: the work involves interpreting language, recognizing patterns, or adapting to context.
- Enough information to do the work: relevant evidence, examples and instructions are available.
- Possible to judge the result: you can test, review or measure whether it worked.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

What is not a good fit for AI automation?

- The rules already solve it: use conventional software for exact calculations and fixed logic.
- The necessary context is missing: AI cannot reliably fill gaps in facts, evidence or business decisions.
- Mistakes cannot be caught before they cause harm: keep AI in an assisting role until suitable checks exist.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The 3R principle™

Real

Is this a problem worth solving?

Repetitive

Does the work recur often enough?

Ready

Can we access the data and systems it needs?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

From individual problems to enterprise transformation

A company has many AI opportunities.
Where do we start?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The vertical slice framework™

1. Choose a business outcome.
2. Understand the workflow.
3. Solve a bounded problem.
4. Connect and expand what works.

One complete slice, then adjacent work

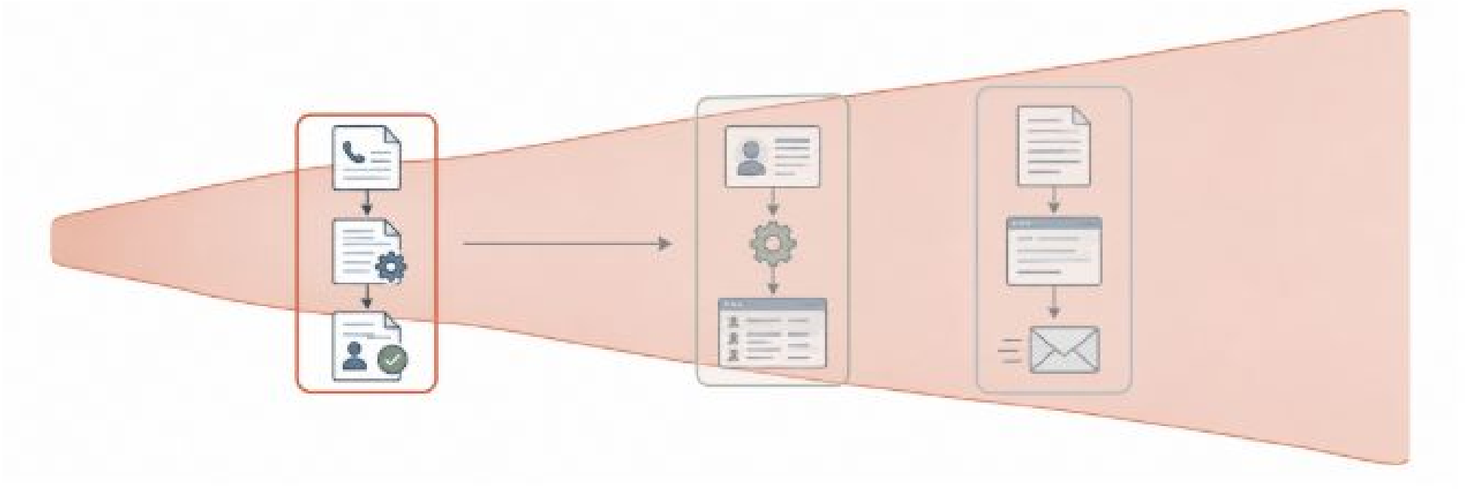

Inputs → work → reviewed result

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Step 1: Choose a business outcome

Get accurate proposals to customers sooner.

Lead generation → Discovery → Proposal → Follow-up → Close

Lead generation works

Next bottleneck: proposals wait

Measure turnaround, corrections and review effort.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Step 2: Understand the workflow

Lead generation → Discovery → Proposal → Follow-up → Close

Inside proposal preparation

Customer needs → Approved pricing → Draft → Review + send

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Step 3: Solve a bounded problem

Scoping the first problem: one team, one standard proposal.

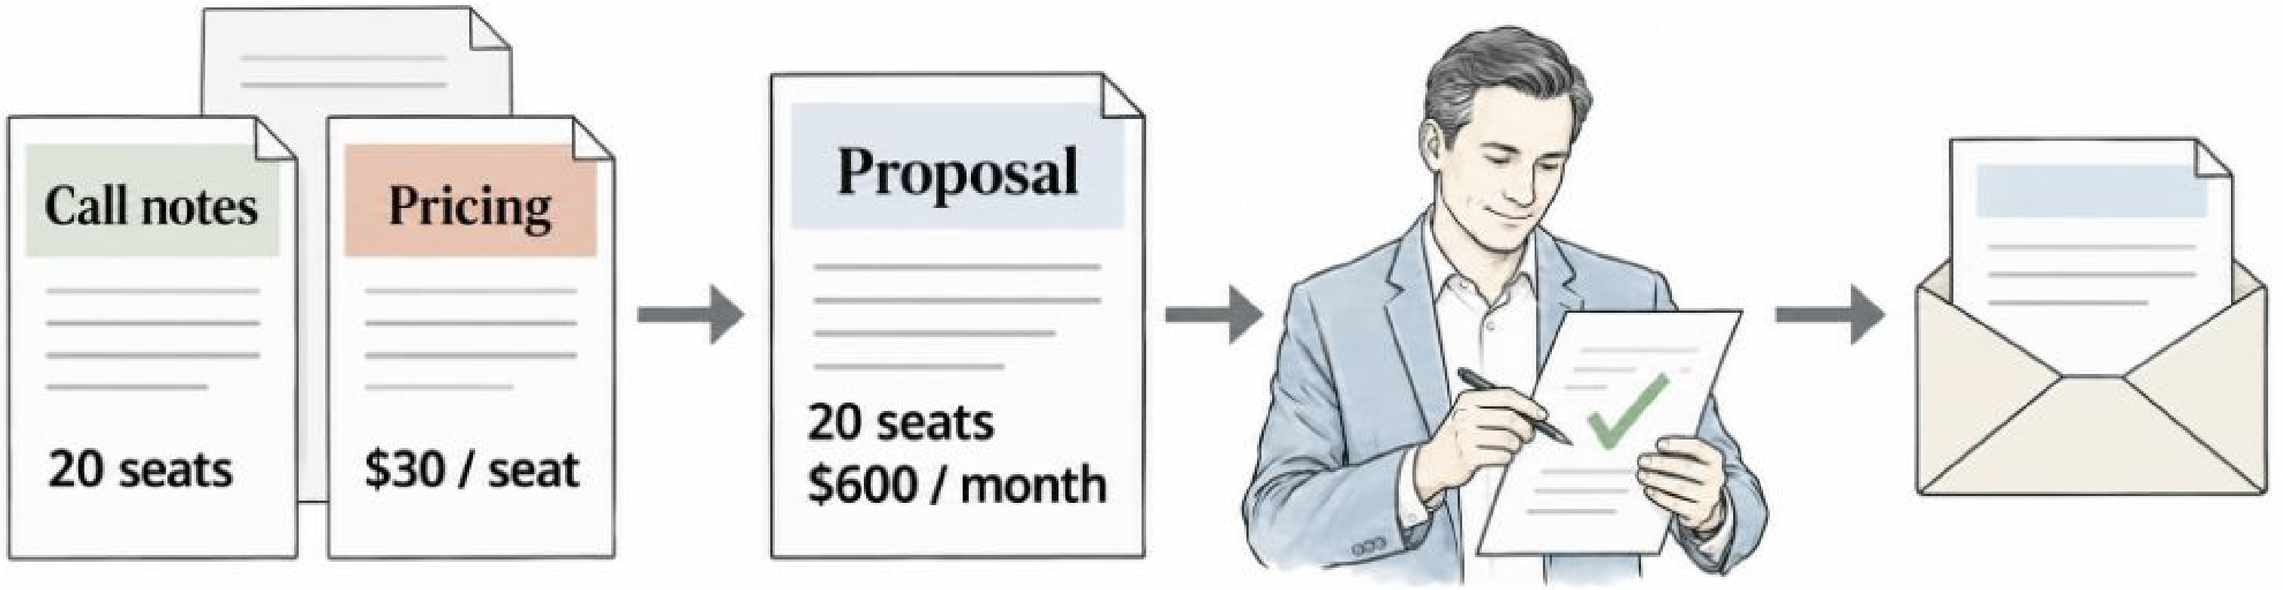

Approved inputs | Draft | Rep checks | Rep sends

Monthly pricing shown; special prices or terms go to a person.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Step 4: Connect and expand what works

Lead generation → Discovery → Proposal → Follow-up → Close

Working

Proposal slice proves useful

Next candidate: follow-up

Check the impact, then choose the next bottleneck.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Expanding across the enterprise

Sales

Leads → proposals → follow-up

Finance

Documents → explanations

Engineering

Bug reports → reviewed fixes

Procurement

Quotes → purchase requests

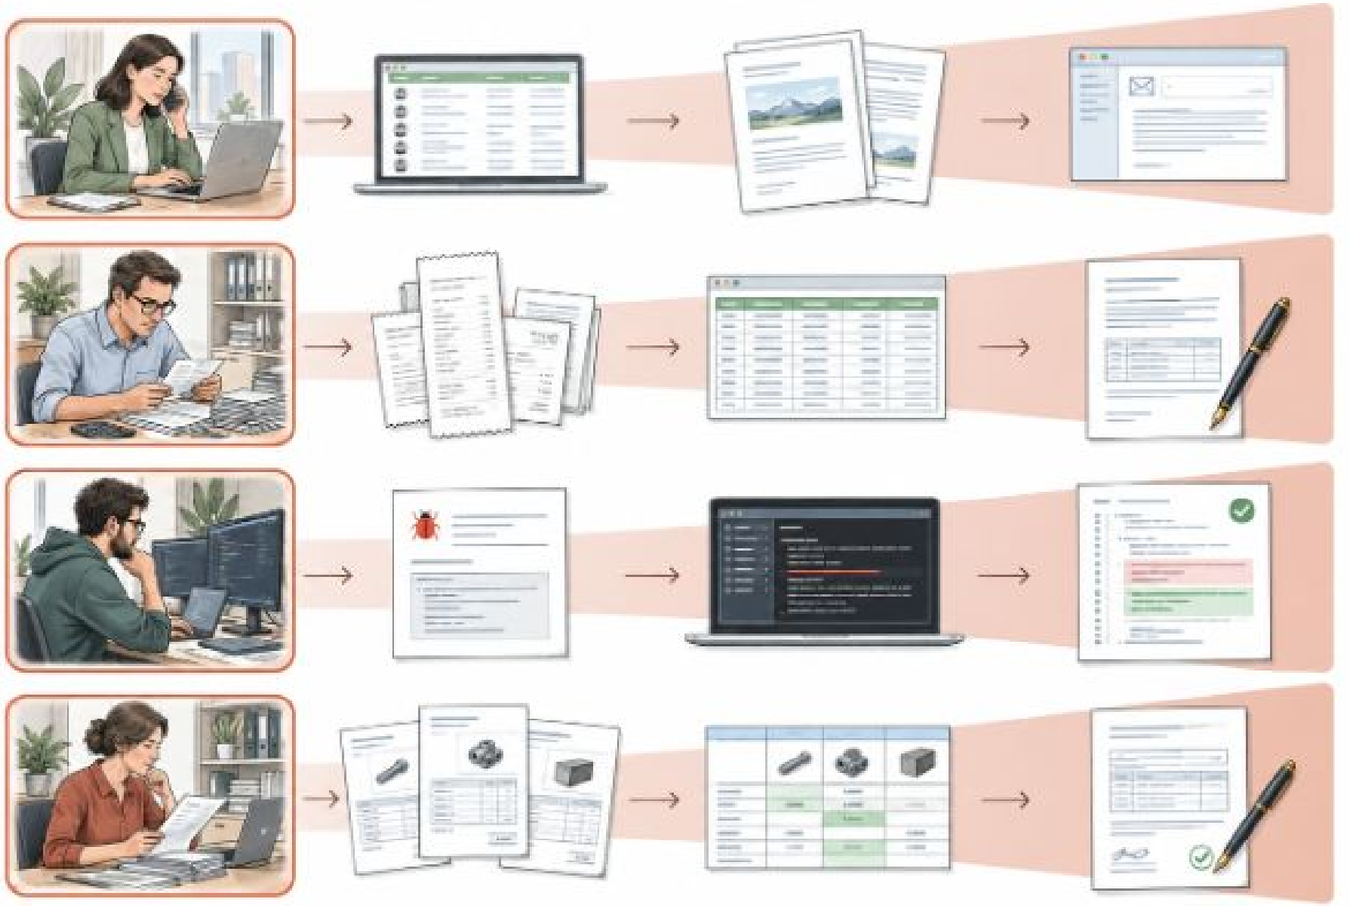

Start with one useful workflow.

Expand where it proves worthwhile.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

LevelUp Lens™

Choose a business outcome, understand the workflow, solve one bounded problem, then use what works to expand across the business.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

04 System Design for AI

Three reasons AI applications need a different approach to system design.

- Non-determinism
- The agency-control trade-off
- Operational issues

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Components of agentic AI

Add components that let the system act on the model's generations.

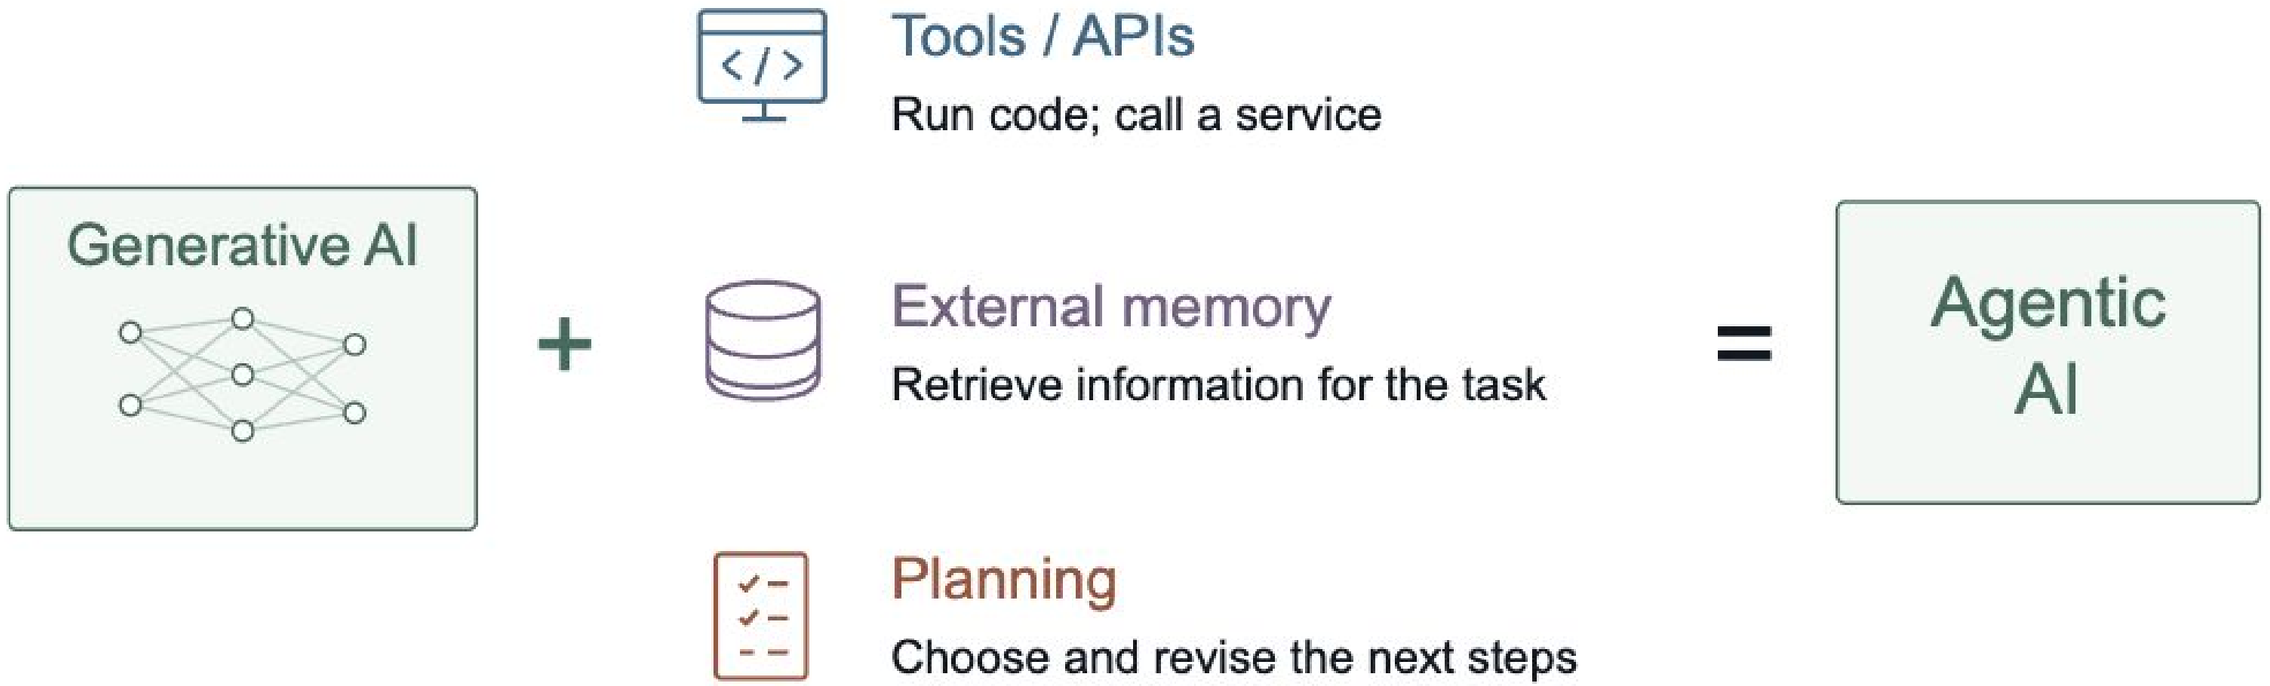

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Non-determinism

How do we build dependable applications when model outputs can vary?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Generative AI is probabilistic

An LLM assigns next-token probabilities across a large vocabulary.

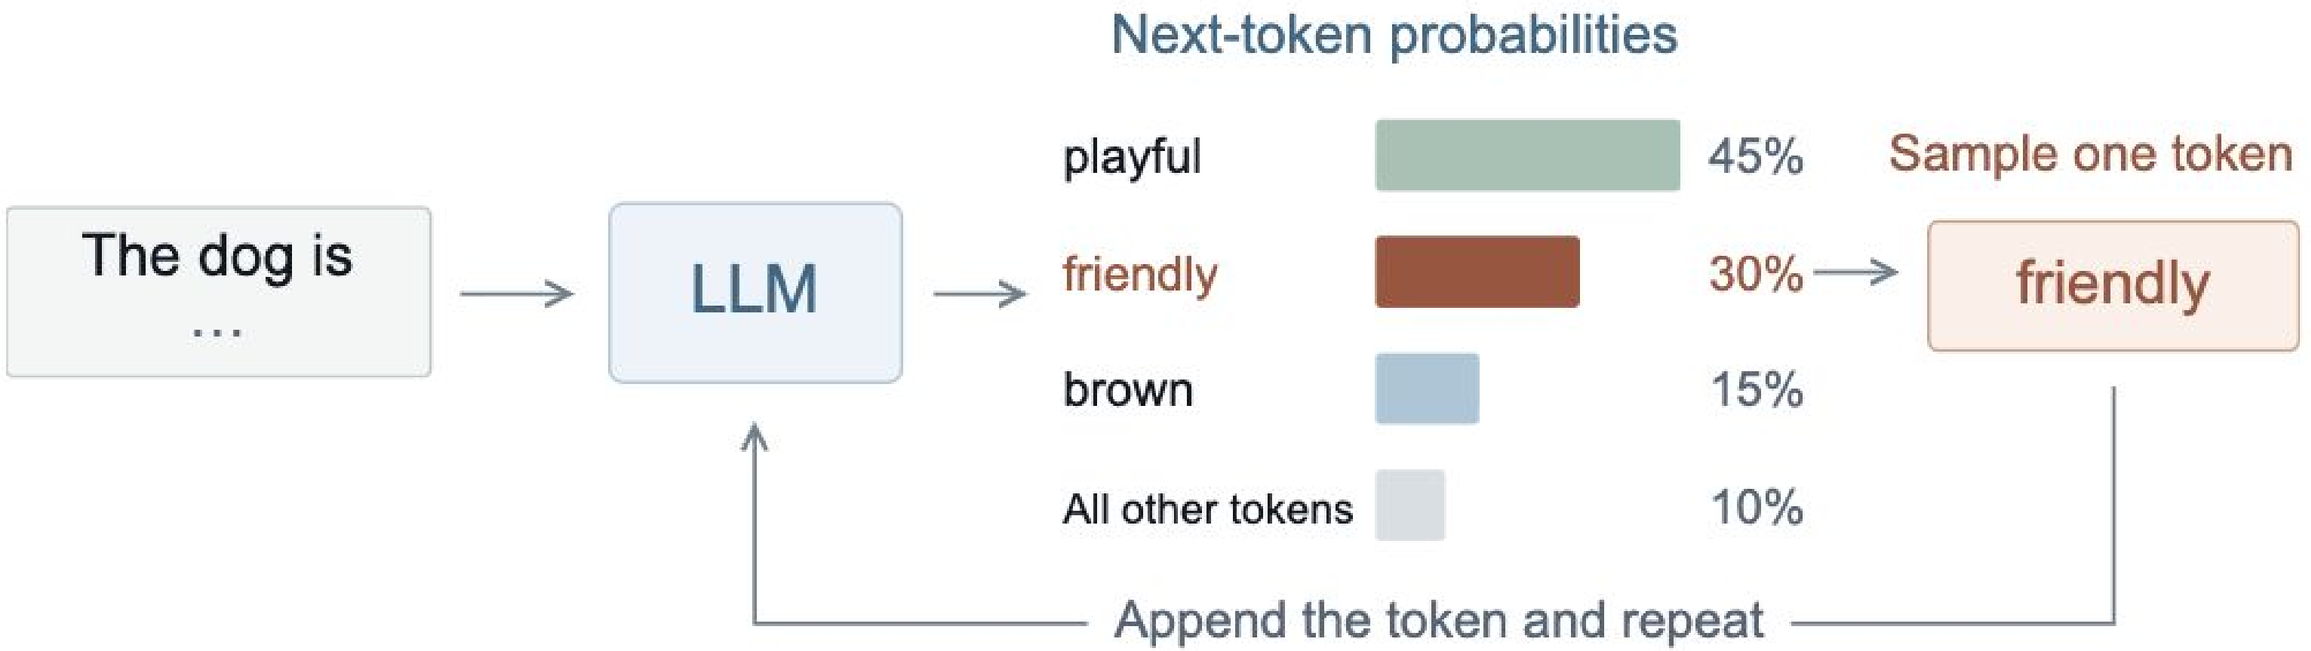

Sampling can make generation non-deterministic: the same prompt can produce different responses.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Programming, predictive AI and generative AI

Traditional programming

Write the rules.

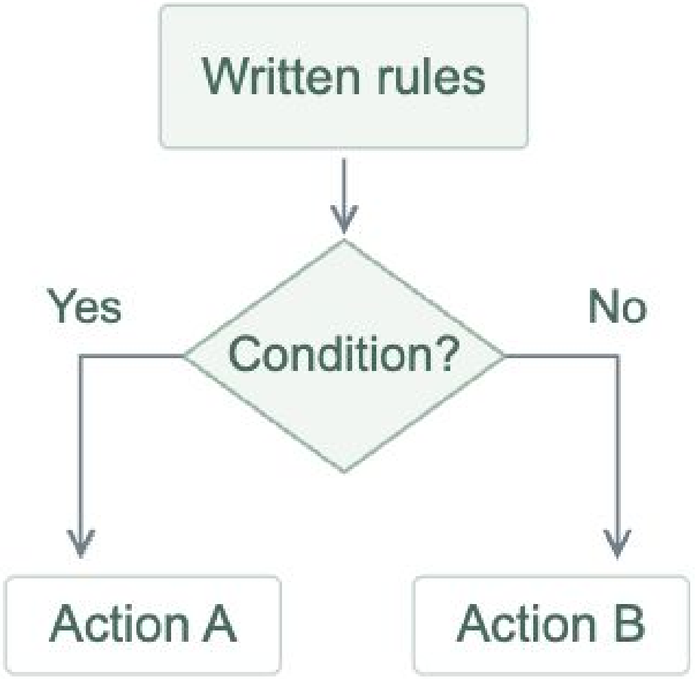

Predictive AI

Learn to predict a result.

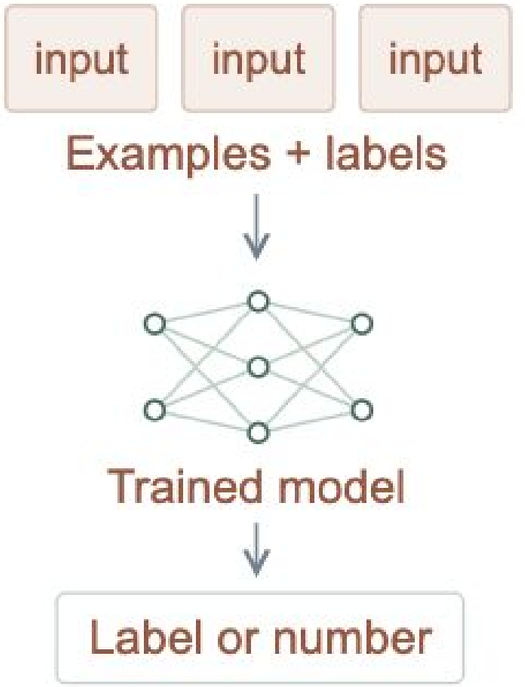

Generative AI

Learn to generate content.

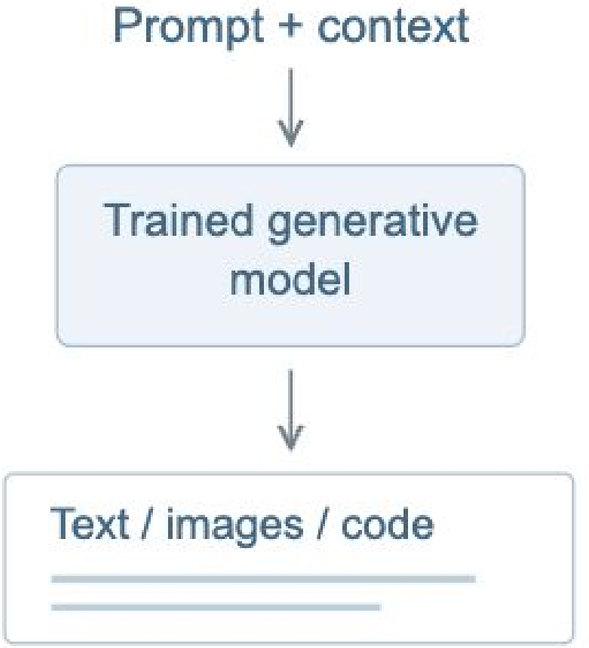

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Image classification: three approaches

Does this image represent a cat or a dog?

Traditional programming

Handwritten feature rules

Predictive AI

Trained on labeled images

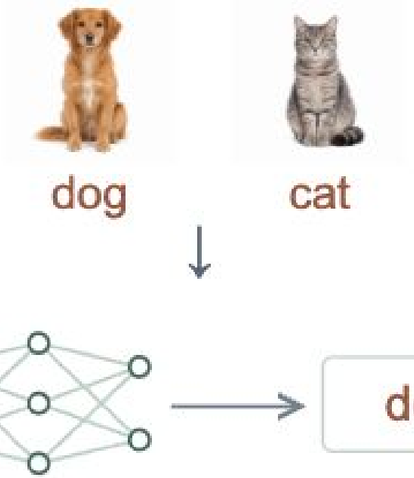

Generative AI

Image + question

↓

"This image shows a brown dog."

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The determinism framework

Deterministic behavior repeats the result under the same inputs and execution conditions.

Approach	Process	Output behavior
Programming	Run specified logic	Repeatable when inputs and state stay fixed
Predictive AI	Apply a trained mapping	A label or number can be repeatable
Generative AI	Generate from a learned distribution	Sampling can vary the wording: dog → a brown dog

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Generative AI: variation in outputs

Does this image represent a cat or a dog?

dog

yes, this is a cute brown dog

well, this isn't a cat for sure, it's a dog

All three communicate "dog." Can the application use all three?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

AI outputs in a recommendation system

A stored label. Recommend dog images from a liked image.

Image A matches the stored label and is included.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

AI outputs in a recommendation system

A sentence as the label. Recommend dog images from a liked image.

Image A

→

Classify

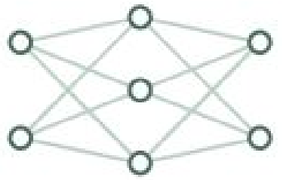

Cat or dog?

→

Store the output

Image	label
A	yes, this is a cute brown dog
B	dog
C	cat

→

Retrieve images

```
WHERE label = 'dog'
```

↓

Shown to the user

B

Image A is missed.

The sentence identifies a dog, but it does not equal the label dog.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

AI outputs in a recommendation system

Another sentence as the label. Recommend dog images from a liked image.

Image A

Classify

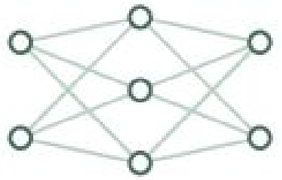

Cat or dog?

Store the output

Image	label
A	well, this isn't a cat for sure, it's a dog
B	dog
C	cat

Retrieve images

```
WHERE label = 'dog'
```

Shown to the user

B

Image A is missed.

The sentence identifies a dog, but it does not equal the label dog.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Enterprises need deterministic outcomes from non-deterministic technology.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Output contracts

Model output

dog

yes, this is a
cute brown dog

→

Allowed-label check

Accept only **cat** or **dog**.

Sentence → reject

→

Store and retrieve

A · dog

```
WHERE label = 'dog'
```

Format check

Is the value in the allowed set?

Task check

Is that label right for this image?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The agency-control trade-off

Which decisions should AI make, and which should remain with people?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

User-driven workflows

To book a hotel, the person makes the choices at each step.

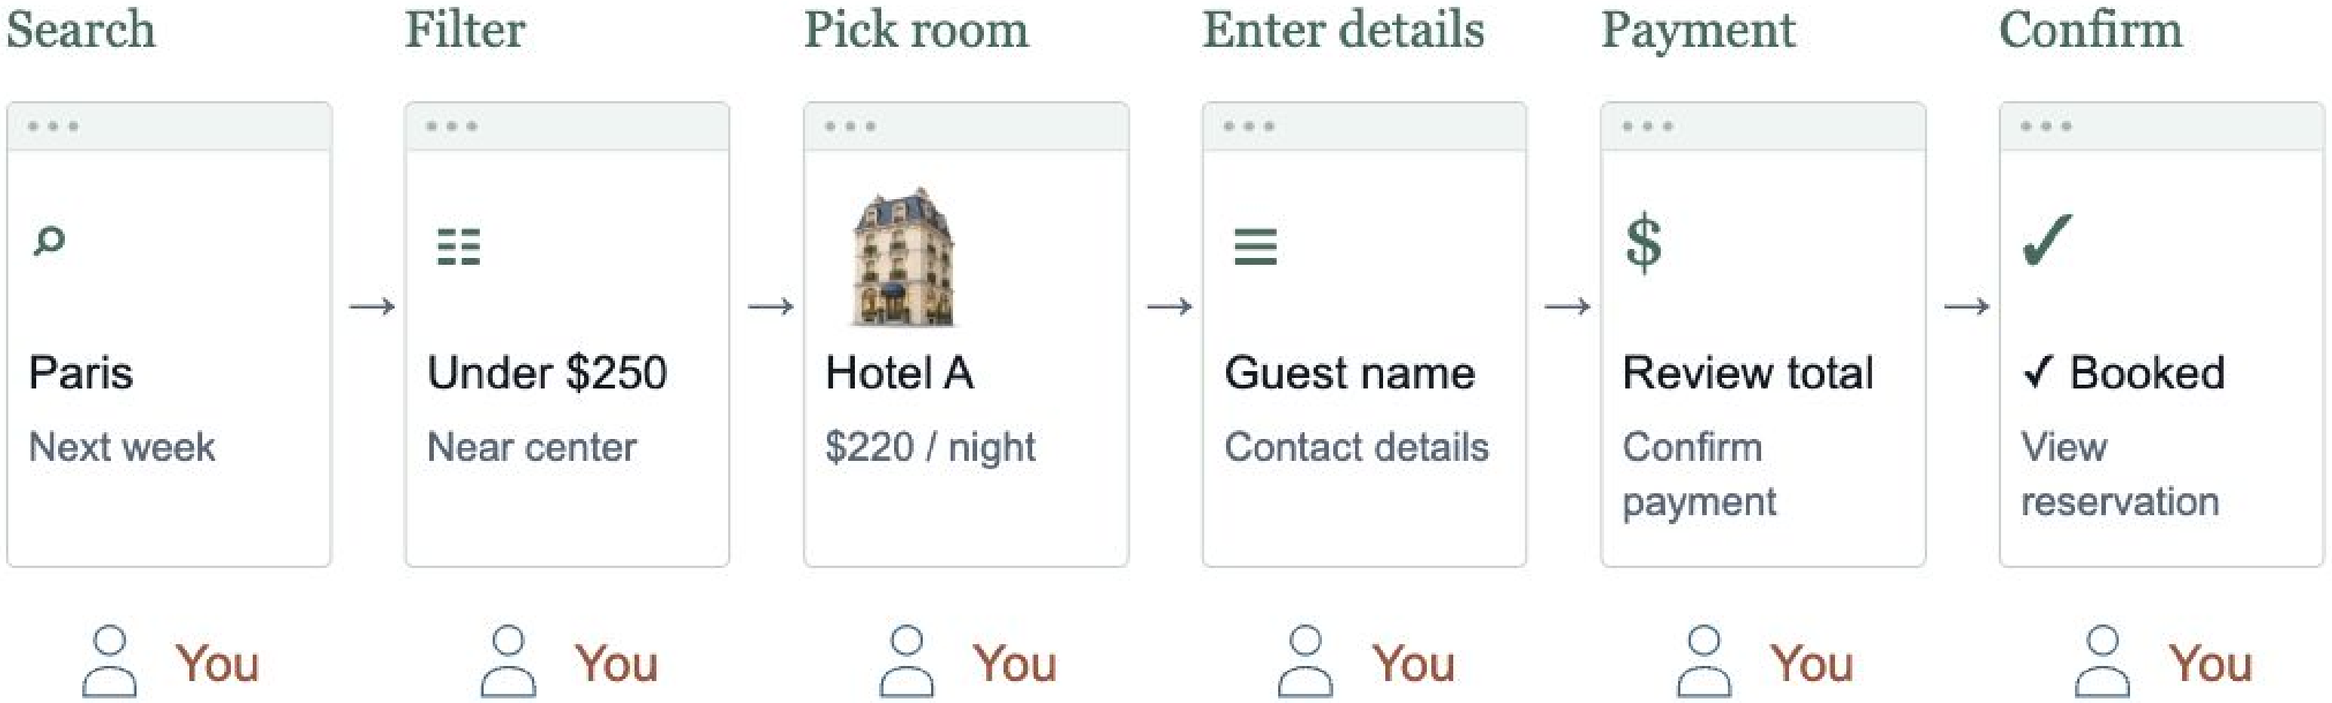

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

User-driven and intent-driven workflows

> "Book a hotel in Paris next week, under $250, near the center."

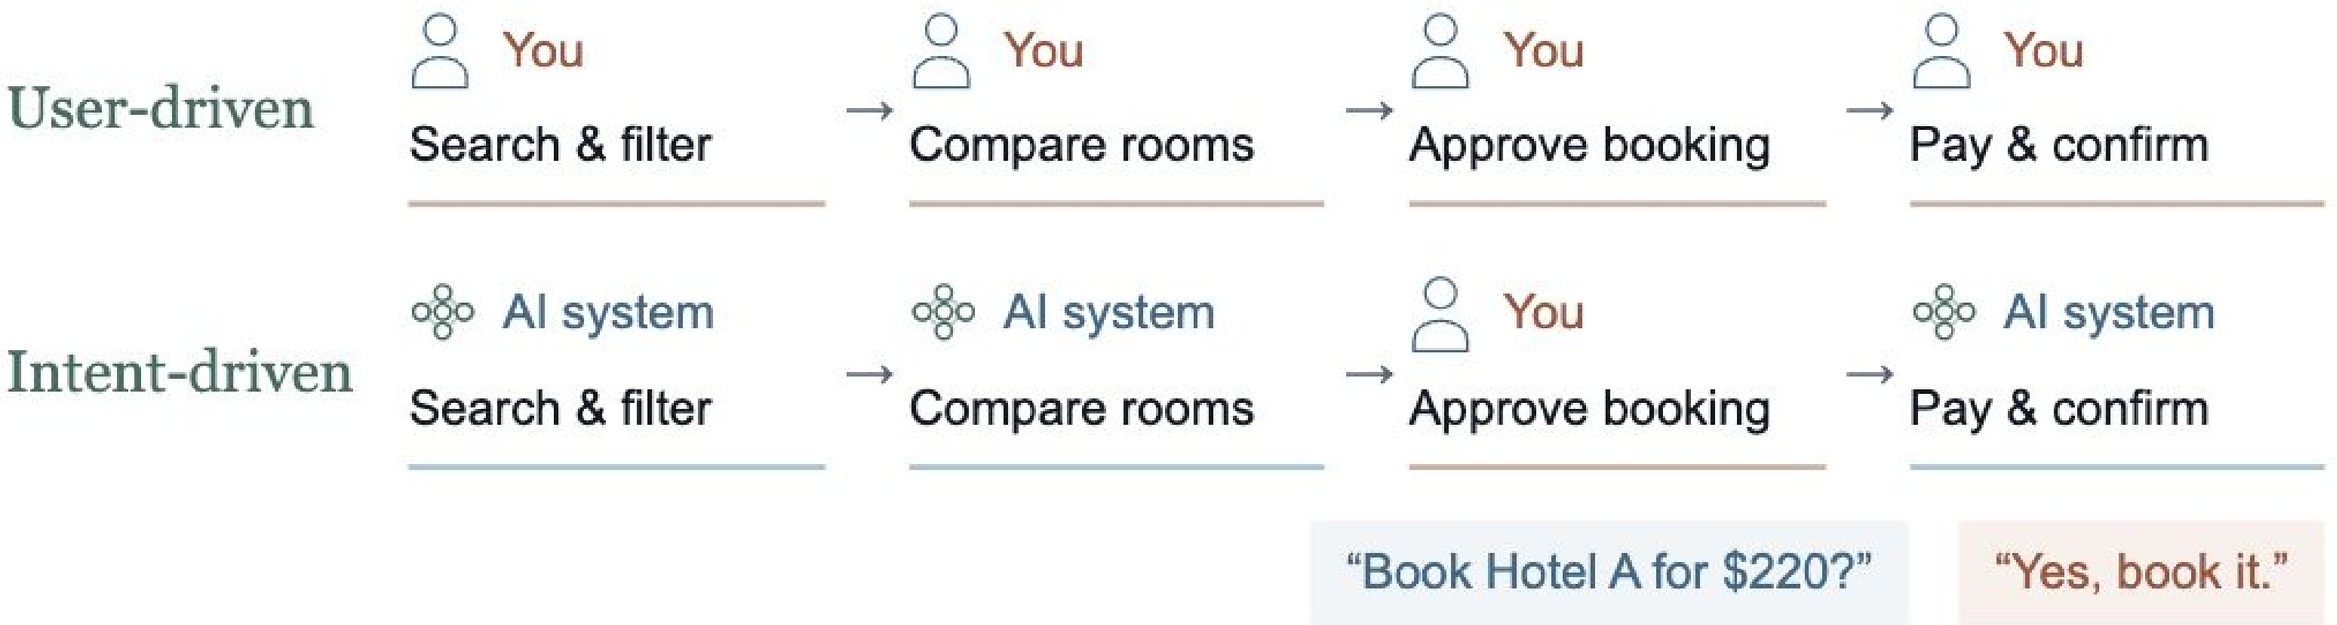

Agency is the authority to take actions toward a goal.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

The agency-control trade-off

Choose which decisions the system may make and which remain with the person.

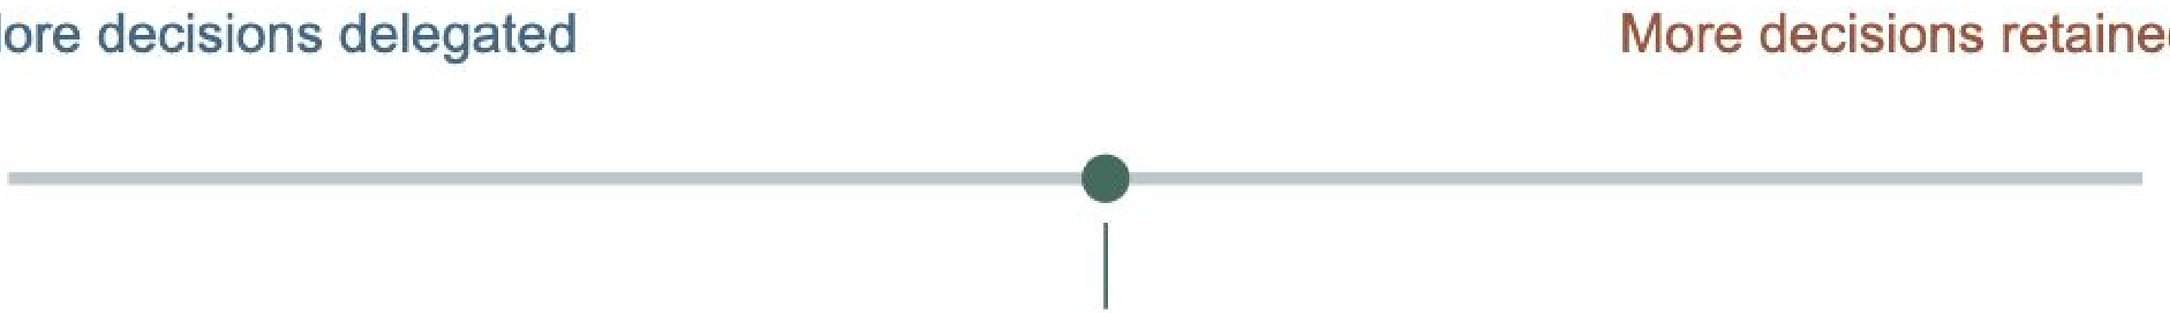

Unresolved requests return to the person.

Delegate only when the system can handle the use case reliably.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Operational issues

What else can make an AI application unreliable?

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.

Operational issues in AI applications

Issue	Explanation
Hallucinations	Produces claims that are false or unsupported.
Sycophancy	Agrees with the user at the expense of accuracy.
Inconsistent response length	Gives too little or too much detail for the task.
Prompt sensitivity	Small wording changes can lead to different results.
Bias reinforcement	Repeats or amplifies biases in data or context.
Limited explainability	Its explanation may not reflect how it reached the answer.
Latency	Model generation and tool calls can delay the result.

© LevelUp Labs. Course materials for enrolled participants only. Do not copy or redistribute.